

Architecture Katas

Winter 2025: Architecture & AI

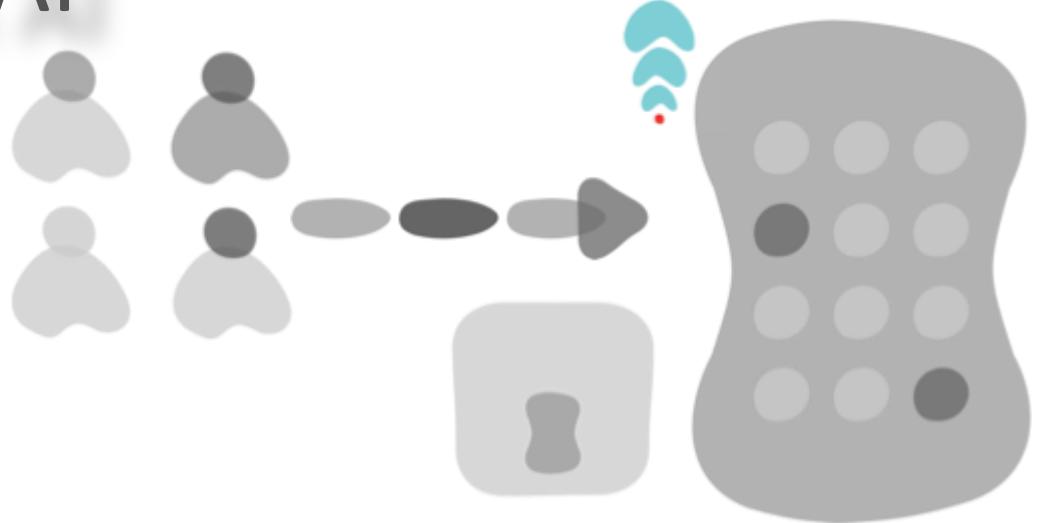


Neal Ford

Thoughtworks

Director / Software Architect / Meme Wrangler

<https://www.nealford.com>



Mark Richards

Independent Consultant

Hands-on Software Architect, Published Author

Founder, DeveloperToArchitect.com

@markrichardssa

Contest Kickoff

Introduction

Meet the Judges

Birgitta Böckeler

Birgitta Böckeler is global lead for AI-assisted software delivery at Thoughtworks and a software developer, architect, and technical leader with 20 years of experience. As a software delivery consultant, Birgitta has had the opportunity to see many organizations and teams succeed, or fail, at delivering valuable software.



Javiera Laso

Javiera Laso is a senior software engineer and architect with a background in biotechnology, specializing in modernizing legacy systems into scalable, efficient, and maintainable solutions. She applies Domain-Driven Design (DDD) principles to create software that meets technical requirements while aligning with business goals. Javiera is experienced in designing modular, cloud-ready architectures, using techniques such as microservices and event-driven design to tackle complex challenges. Committed to continuous learning and improvement, she mentors teams, encourages best practices in software development, and ensures that her solutions are both practical and future-proof. Her approach blends technical expertise with a clear understanding of business objectives, making her a valuable asset in any software development project.



Andrew Harmel-Law

Andrew Harmel-Law is an overenthusiastic tech principal for Thoughtworks, where he specializes in domain-driven design, Java/JVM technologies, Agile delivery, build tools and automation, and organization design. He's motivated by the efficient delivery of large-scale software solutions, and he understands that people, architecture, process, and tooling all have key roles to play in achieving this. His experience spans the software development lifecycle across many sectors including government, banking, and ecommerce.

Andrew also has a passion for open source software and its communities, and he enjoys open-sourcing his code. He's been involved with OSS as a user, contributor, expert group member, and paid advocate—most notably as one of the Jenkins Job DSL originators. He shares his experiences through consulting, mentoring, writing blog posts, and speaking at and organizing conferences.



Where did this
idea come from?

The screenshot shows a web browser window with the URL `archkatas.herokuapp.com` in the address bar. The page title is "Architectural Katas". Below the title is a quote by Fred Brooks: "How do we get great designers? Great designers design, of course." --Fred Brooks. Underneath is another quote by Ted Neward: "So how are we supposed to get great architects, if they only get the chance to architect fewer than a half-dozen times in their career?" -Ted Neward. A blue button labeled "Do one! »" is visible. The main content area has several sections: "About", "Rules", "Contribute", "Invite", "Lead", and "Join". Each section has a brief description and a call-to-action button.

Architectural Katas

"How do we get great designers? Great designers design, of course." --Fred Brooks

"So how are we supposed to get great architects, if they only get the chance to architect fewer than a half-dozen times in their career?"
-Ted Neward

[Do one! »](#)

About
The Architectural Katas started as a presentation workshop by Ted Neward. They've taken on a life of their own.
[Learn more »](#)

Invite
Want an experienced Architectural Kata moderator to run the workshop at your place of business?
[Contact »](#)

Rules
Doing an Architectural Kata requires you to obey a few rules in order to get the maximum out of the activity.
[Read rules »](#)

Lead
Want to run the Architectural Katas yourself? There's only a few things you need to know before you do.
[Learn how »](#)

Contribute
New Kata problems/proposals are always welcome.
[Send ideas »](#)

Join
Want to find a group near you that's running the Architectural Katas?
[Find groups »](#)

© Neward & Associates 2012

Architectural Katas

"How do we get great designers? Great designers design, of course." --Fred Brooks

"So how are we supposed to get great architects, if they only get the chance to architect fewer than a half dozen times in their career?"

--Ted Neward

[Do one! »](#)

Where did this idea come from?

About

The Architectural Katas started as a presentation workshop by Ted Neward. They've taken on a life of their own.

[Learn more »](#)

Invite

Want an experienced Architectural Kata moderator to run the workshop at your place of business?

Rules

Doing an Architectural Kata requires you to obey a few rules in order to get the maximum out of the activity.

[Read rules »](#)

Lead

Want to run the Architectural Katas yourself? There's only a few things you need to know before you do.

Contribute

New Kata problems/proposals are always welcome.

[Send ideas »](#)

Join

Want to find a group near you that's running the Architectural Katas?

• • and then • •

<http://fundamentalsofsoftwarearchitecture.com/katas/>

The screenshot shows a web browser window with the following details:

- Address Bar:** Not Secure — fundamentalsofsoftwarearchitecture.com
- Page Title:** fundamentalsofsoftwarearchitecture.com
- Page Content:**
 - Navigation links: Architectural Katas, Updated Fundamentals of Software Architecture Images, Architectural Katas, Fundamentals of Software Architecture, List of Architecture Katas.
 - Section Header:** Architectural Katas
 - Text:** *inspired by Ted Neward's original [Architectural Katas](#)*
 - Text Block:** "How do we get great designers?
Great designers design,
of course."
— Fred Brooks
 - Text Block:** "So how are we supposed to get great architects, if
they only get the chance to architect fewer than
a half-dozen times in their career?"
— Ted Neward
 - Section Header:** About
 - Text:** Architectural Katas are intended as a small-group (3-5 people) exercise, usually as part of a larger group (4-10 groups are ideal), each of whom is doing a different kata. A Moderator keeps track of time, assigns Katas (or allows this website to choose one randomly), and acts as the facilitator for the exercise.
 - Text:** Each group is given a project (in many ways, an RFP—Request For Proposal) that needs development. The project team meets for a while, discovers requirements that aren't in the original proposal by



...and then...

Architecture Katas

Winter 2025: Architecture & AI

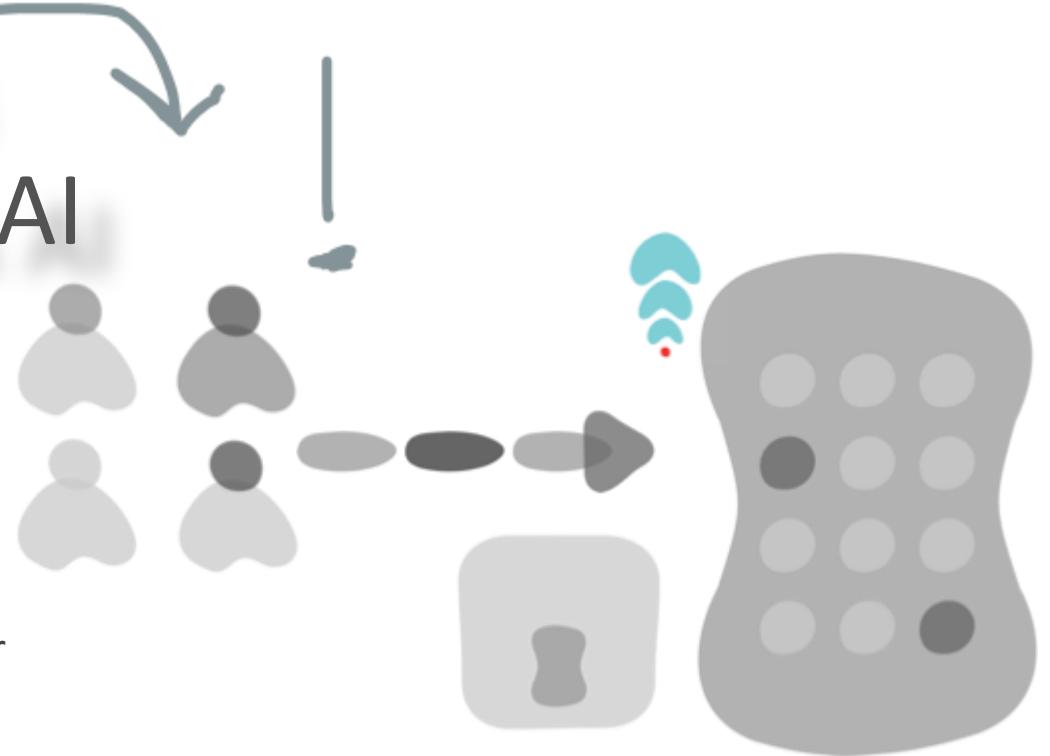


Neal Ford

Thoughtworks

Director / Software Architect / Meme Wrangler

<https://www.nealford.com>



Mark Richards

Independent Consultant

Hands-on Software Architect, Published Author

Founder, DeveloperToArchitect.com

@markrichardssa

Contest Kickoff

Goals:

- Learn about how AI can be applied
- Learn about the tools and categories within the AI ecosystem
- Learn about AI architectural patterns and how they can be applied
- Learn about the limitations and risks associated with AI in architecture
- Learn about how AI impacts the structure of a system

A Historical Perspective...

We knew it would eventually happen, and sure enough it did. IT Professionals wishing to become a software architect (and existing software architects) must be licensed professions in the same way doctors, nurses, building architects, and lawyers are licensed. This law was passed in the U.S. several years ago, and a Software Architecture Licensing Board (the SALB) formed to accredit companies that administer certification exams. A handful of accredited companies, including Certifiable, Inc., sprang into existence to certify and license new and existing software architects.

Having seen the success of this new law, U.K, Europe, and Asia recently passed similar laws requiring software architecture certification to get or maintain a job as a Software Architect. The licensing boards from these continents all agreed that rather than trying to establish new certification tests and form new certification companies, they would for the time being leverage existing accredited U.S. companies to administer certification tests and license new and existing software architects.



Certifiable, Inc

Additional Context

- The company currently employs 300 expert software architect consultants across the U.S. Of these, 5 are considered “designated” experts who have the authorization to modify certification tests and change/add case studies based on their own analysis and feedback from other employed expert software architects. Expert software architects hired by Certifiable, Inc. earn \$50 per hour for grading certification exams. While this rate varies from company to company, it's in line with what other accredited licensing companies pay for their experts.
- Currently, Certifiable, Inc. has on average 200 candidates per week seeking certification across the U.S. That number is expected to grow 5-10X based on oversees expansion as well as the anticipated 21% growth over the next 4 years.
- The certification exam to obtain software architecture licensing costs \$800.00. This price is fixed and established by the SALB. Overseas nations have agreed to use this cost structure.



Certifiable, Inc

Certifiable, Inc. is an accredited leader in software architecture certification based in the United States. They currently own a large part of the marketshare in the software architecture certification space. Their existing flagship system, **SoftArchCert**, provides accredited certification to qualified software architects. The company currently employs dozens of expert software architects to create and maintain certification tests and review and grade software architecture submissions as part of the certification program.

Recent acceptance of software architecture certification in other parts of the world, including Europe, U.K., and Asia, has put Certifiable, Inc. in a tailspin. Based on the expected increased volume of certification requests (5-10X the normal request volume), the company is worried that the manual processes they currently have in place will not be able to handle the increased demand for software architecture certifications. Therefore, the company would like to see how Generative AI might be able to be applied to handle the increase in demand they are anticipating.

Your job as the architecture team for the *SoftArchCert* system is to determine how Generative AI might be used within the current system to address these rapid and significant growth problems. The company would like your team to identify opportunities for the use of AI in the existing system and redesign the architecture to support those changes.



Certifiable, Inc

Certification Process

- The current certification process requires candidates to pass two different tests: an aptitude test and an architecture submission.
- The first test (aptitude test) involves multiple choice and short answer questions that the candidate takes through a web-based interface. Tests are timed to avoid candidates researching and looking up answers. Once the candidate has completed the test (or the time is up), the candidate answers are stored in a database for later grading. Multiple choice questions are auto-graded, whereas the short answer questions are graded by employed expert software architects (1 week turnaround). A grade of 80% is required to move onto the second test, which the details and links are emailed to the candidate once graded.
- The second test is a case study that the candidate must create an architecture for. There are five possible case studies available that are assigned at random for each candidate. The candidate has two weeks from the time they download the case study to the time they upload their architectural solution. Once uploaded, the employed expert software architects review the architecture submission and assign a grade based on a specific set of criteria. The expert software architect graders have 1 week to review the architecture submission once it is submitted by the candidate. The candidate submission are assigned randomly to the expert software architects. A comprehensive score of 80% is required to pass the second test and receive the accredited certification.
- If the candidate passes both tests, they receive the accredited certification, and their certification information stored in a certification database which companies can use to verify a candidate's certification and score. Only the candidates name, certification id, date of certification, and score are stored in the database. Companies wishing to verify a candidate's certification can enter this information into the *SoftArchCert verification* website (this information is provided by the candidate). Candidates that fail either test are provided a detailed email from the expert software architect describing why they did not pass the test. The candidate may then reapply to take the certification (candidates who passed the first test only need to retake the architecture submission test).



Certifiable, Inc

Administrative Process

- In addition to grading short answers and architecture submissions, employed Expert Software Architects are responsible for continually analyzing reports and results for the purposes of modifying existing certification tests. For example, if 95% of all candidates miss the same question, analysis should be done to modify the question or remove it. Also, as new techniques, practices, and patterns emerge in the industry, expert software architects add questions to reflect these advances.
- Designated employed expert software architects are also responsible for periodically editing and creating new case studies for the second test to prevent case studies and answers from leaking out to the internet.
- Employed expert software architects are contractors (freelance workers) who work on an hourly basis with contracts for time they work per week. The number of hours per week may vary based on demand.
- An administrator from Certifiable, Inc. maintains the expert software architects in the system in terms of profiles and sign-in credentials (experts can also modify their profiles and sign-in credentials once they are added to the system by the administrator).



Certifiable, Inc

Additional Context

- Currently, there are over 176,000 software architects in the US, with 300,000 job openings yet to be filled. That number across the U.K., Europe and Asia is estimated at 600,000. This field is expected to grow by 21% in the next four years, highlighting the high demand and the challenge of finding qualified professionals (<https://bit.ly/4jBdnDU>).
- Certifiable, Inc. Currently has 120,000 certified software architects in its certification database.
- Certifiable, Inc. guarantees a 1-week turnaround for grading test 1 (aptitude), and a 1-week turnaround for grading test 2 (architecture submission).
- Upon successful completion of test 1, candidates have 30 days to start the architecture submission (test 2), and 2 weeks to complete the architecture submission once started. Downloading the case study signifies the start of test 2.
- Although there are a handful of accredited software architect certification companies and boards, Certifiable, Inc. is recognized as the market leader, with over 80% of companies across the U.S. accepting (or insisting) on Certifiable, Inc's certification. The same is true for overseas expansion.
- Certifiable, Inc. has heard that the use of AI can be somewhat costly and are concerned about significant cost overruns by introducing AI into their certification process. However, while cost is certainly a factor, given that this initiative is a strategic one, they are willing to be somewhat flexible.



Certifiable, Inc

Additional Context

- On average it takes an employed expert software architect 3 hours to grade the short answer questions in Test 1 (aptitude) for a given candidate and provide detailed reasons why an answer was incorrect or partially incorrect (candidate feedback).
- On average it takes an employed expert software architect 8 hours to grade the architecture submission in Test 2 (architecture solution) for a given candidate and provide a detailed analysis of the grading rationale (candidate feedback).



Certifiable, Inc

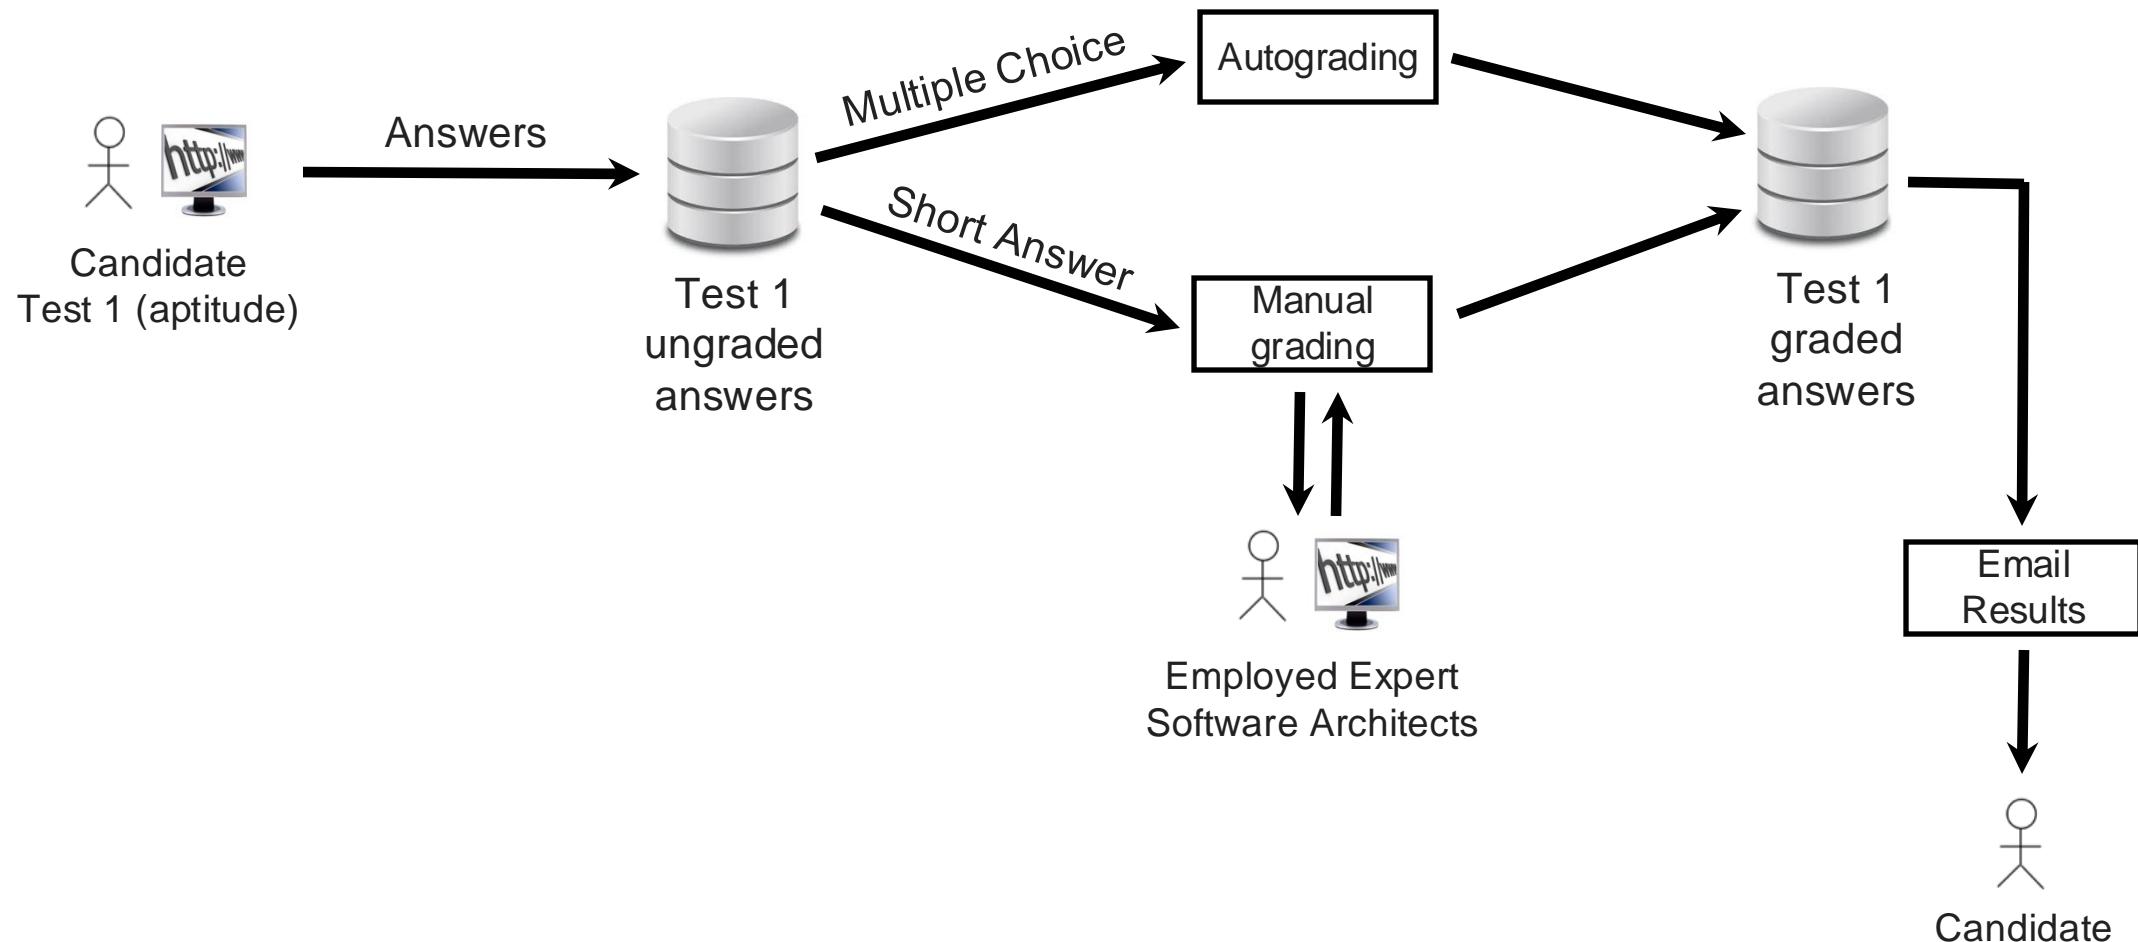
Critical Information

- As a recognized leader in certification, accuracy of tests, case studies, and grading is job and. Inaccurate grading can result in a candidate not getting or maintaining a job and can impact a candidate's career.
- Inaccurate or misleading certification exams and case studies can undermine the credibility of the company's current standing in the marketplace, so accuracy of the certification process is vital for the success of the company.



Certifiable, Inc

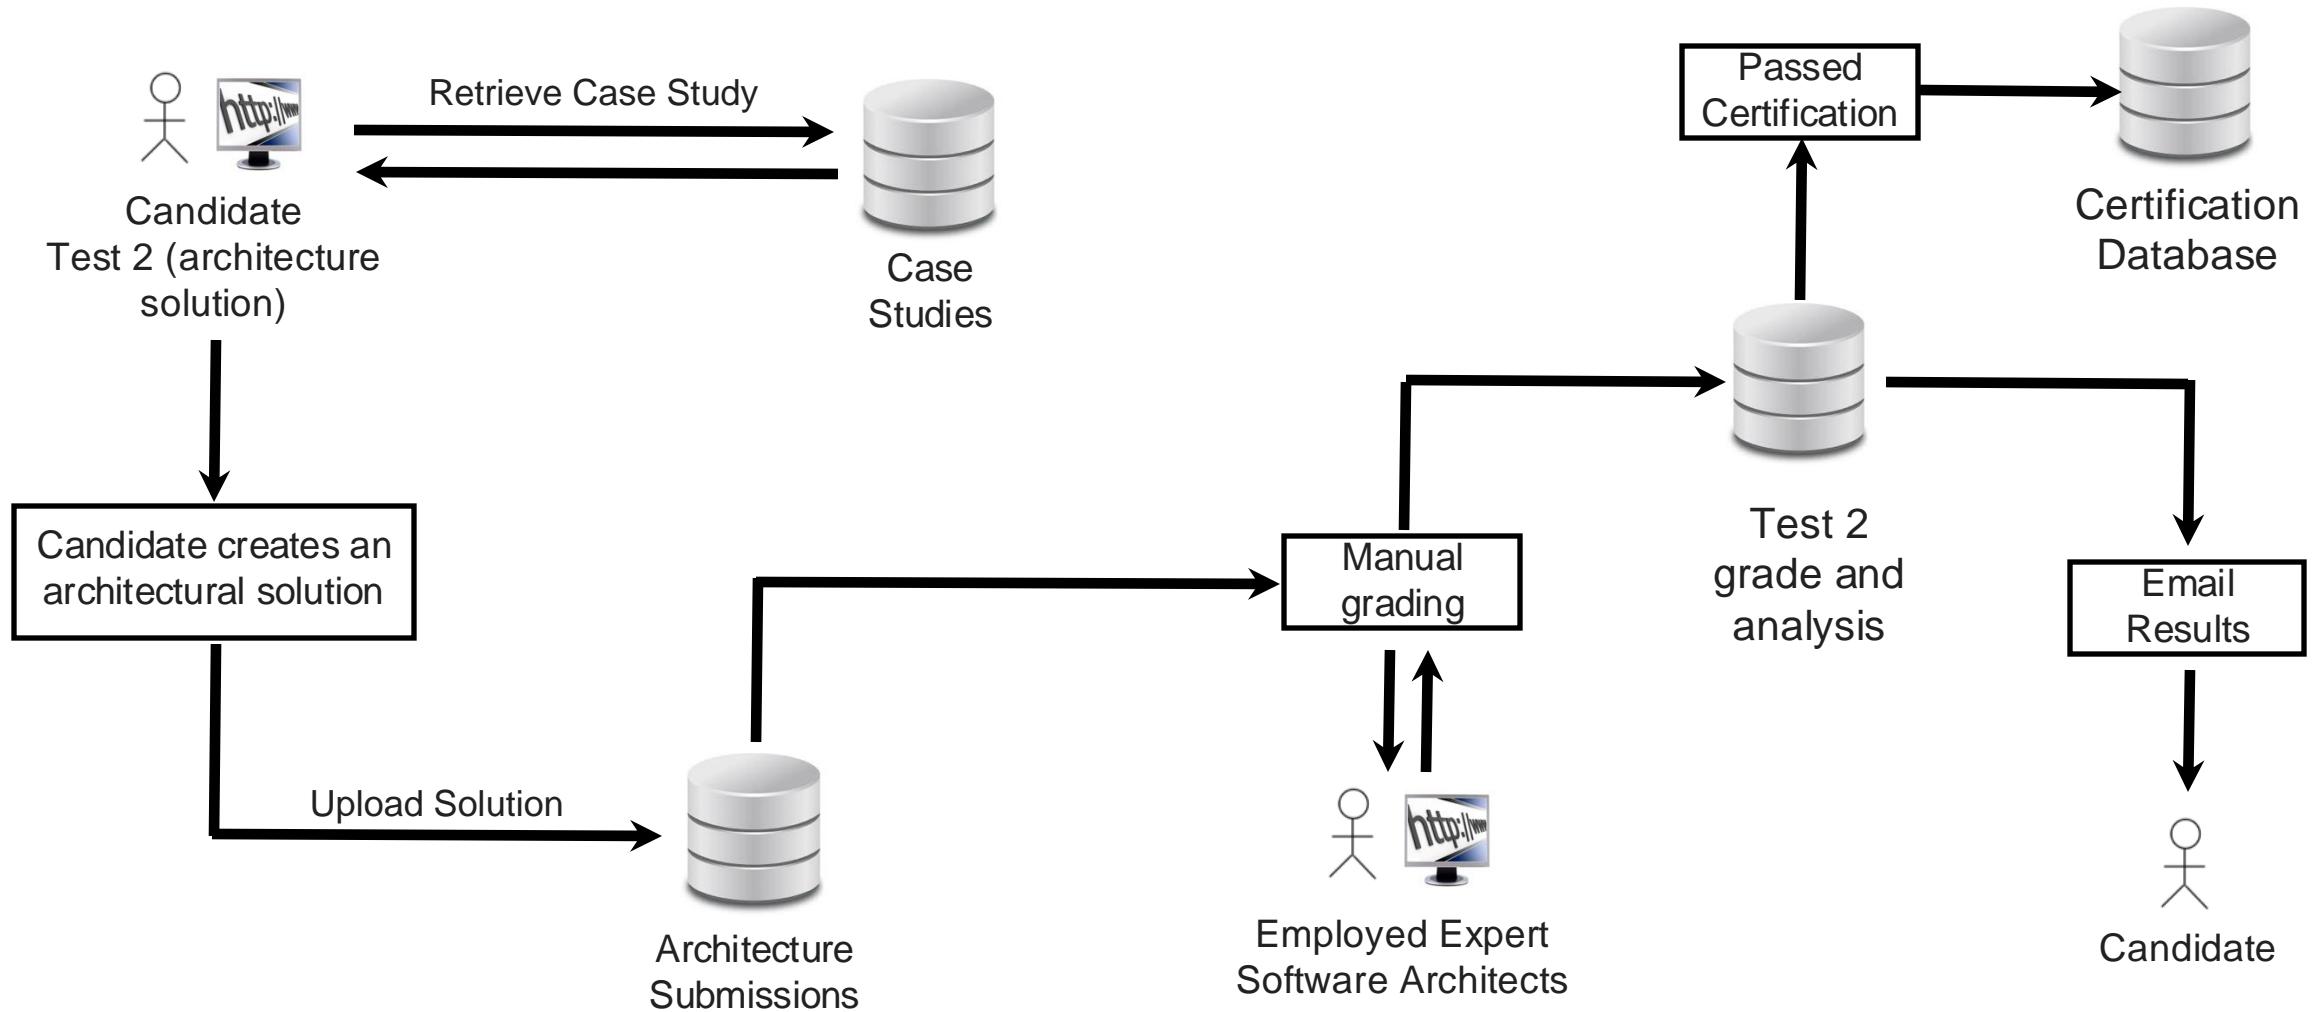
Certification Process (Test 1: Aptitude Test)





Certifiable, Inc

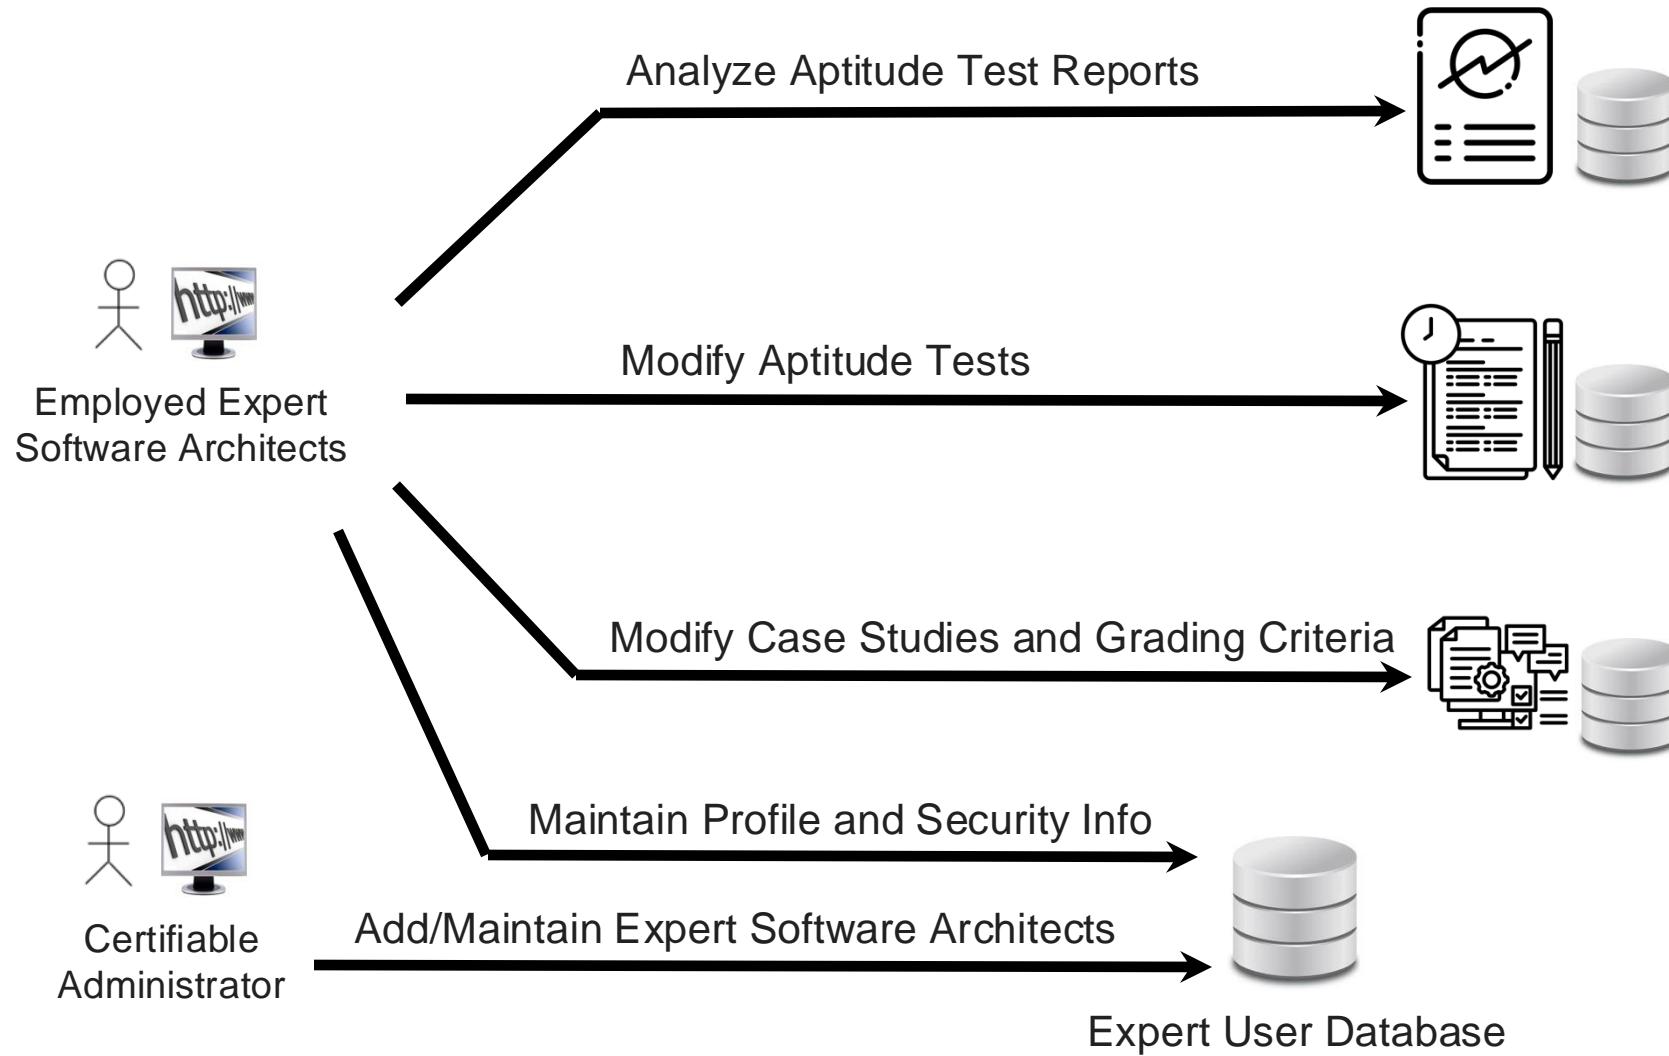
Certification Process (Test 2: Architecture Solution)





Certifiable, Inc

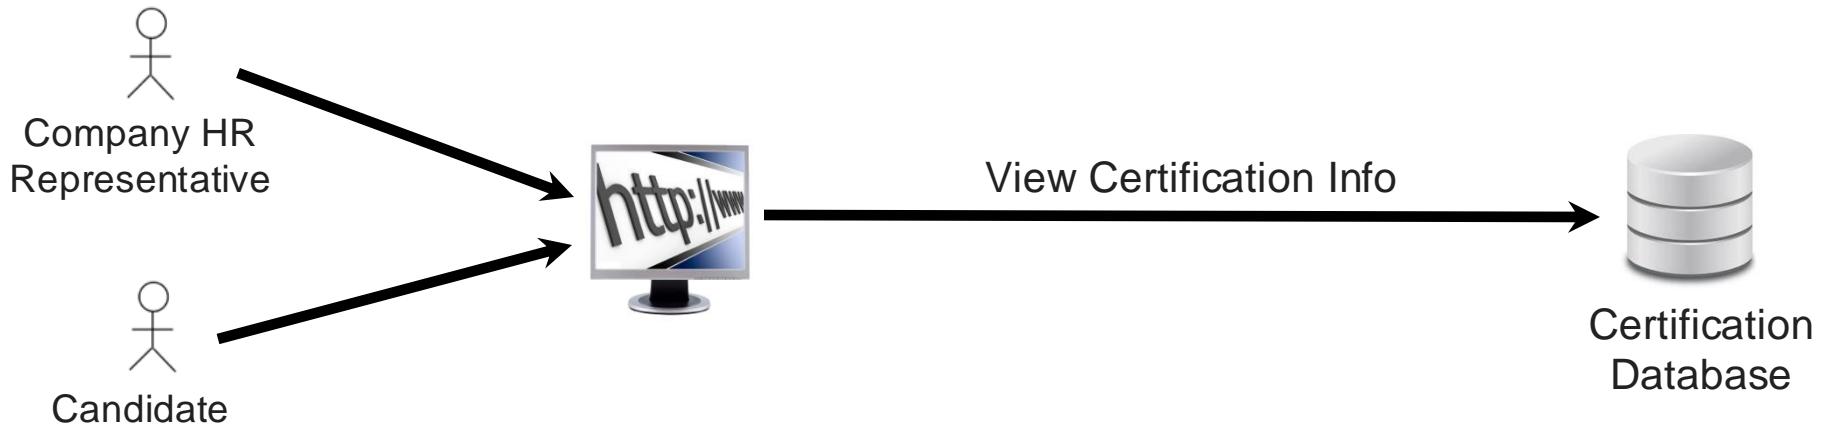
Administration Process



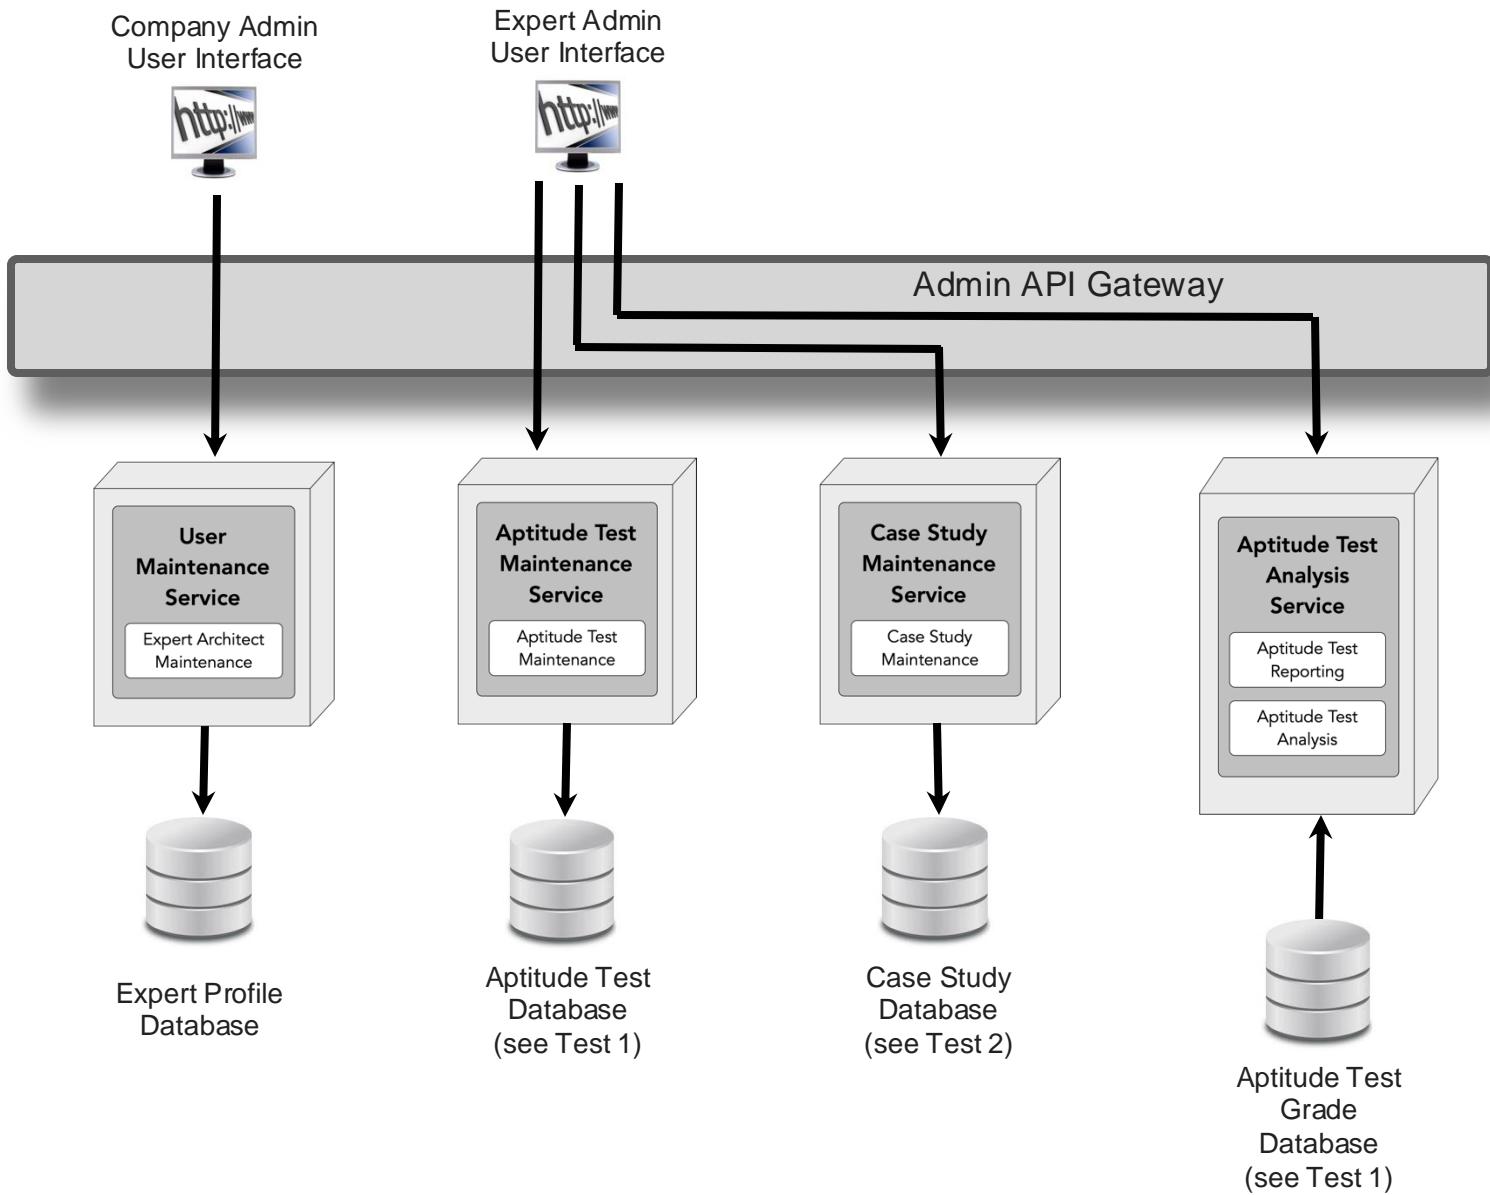


Certifiable, Inc

View Certification Results Process

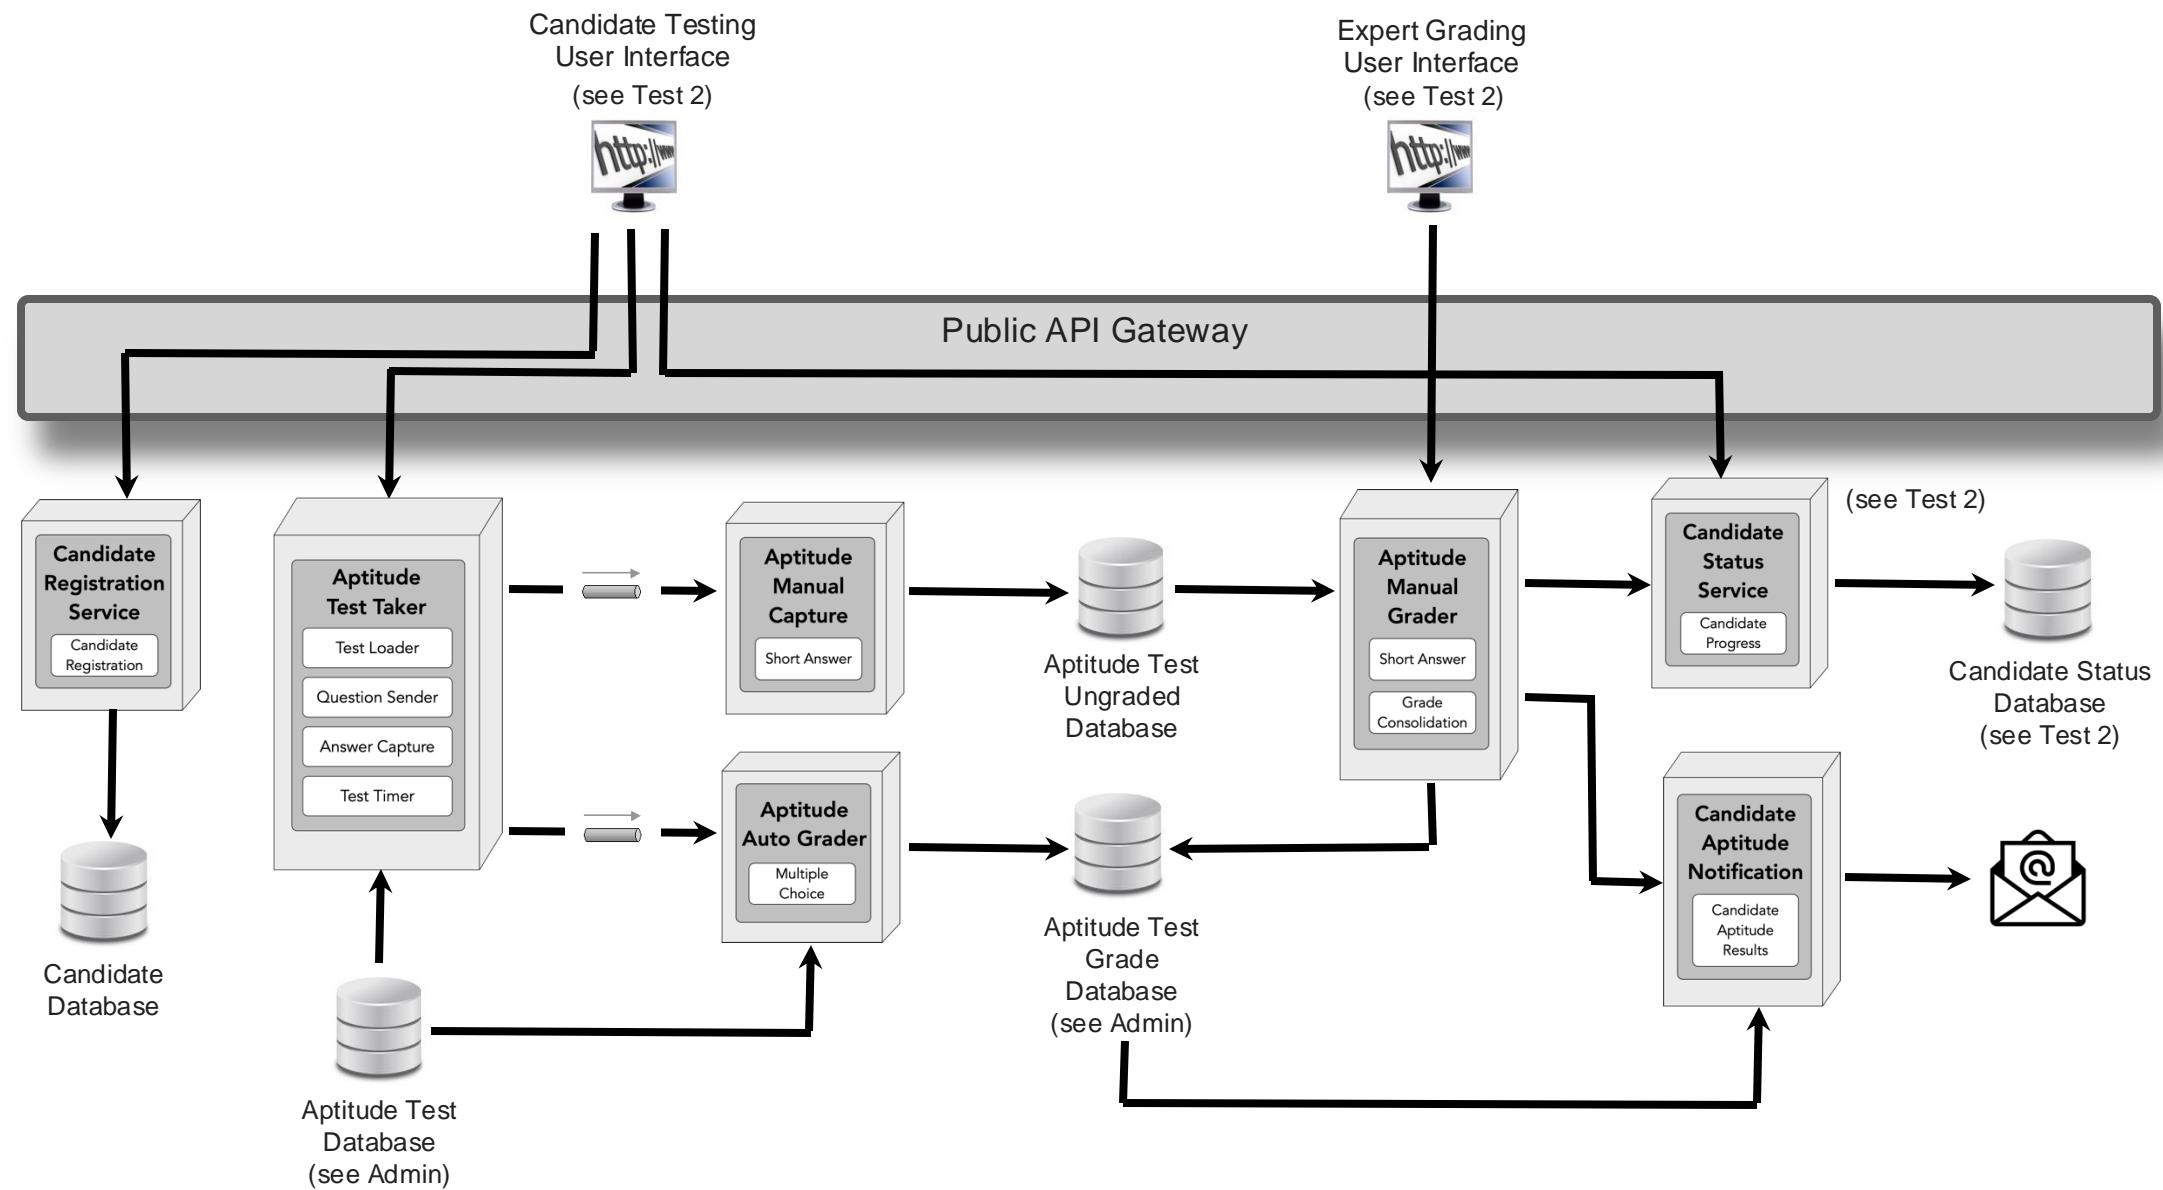


Existing Administrative Architecture



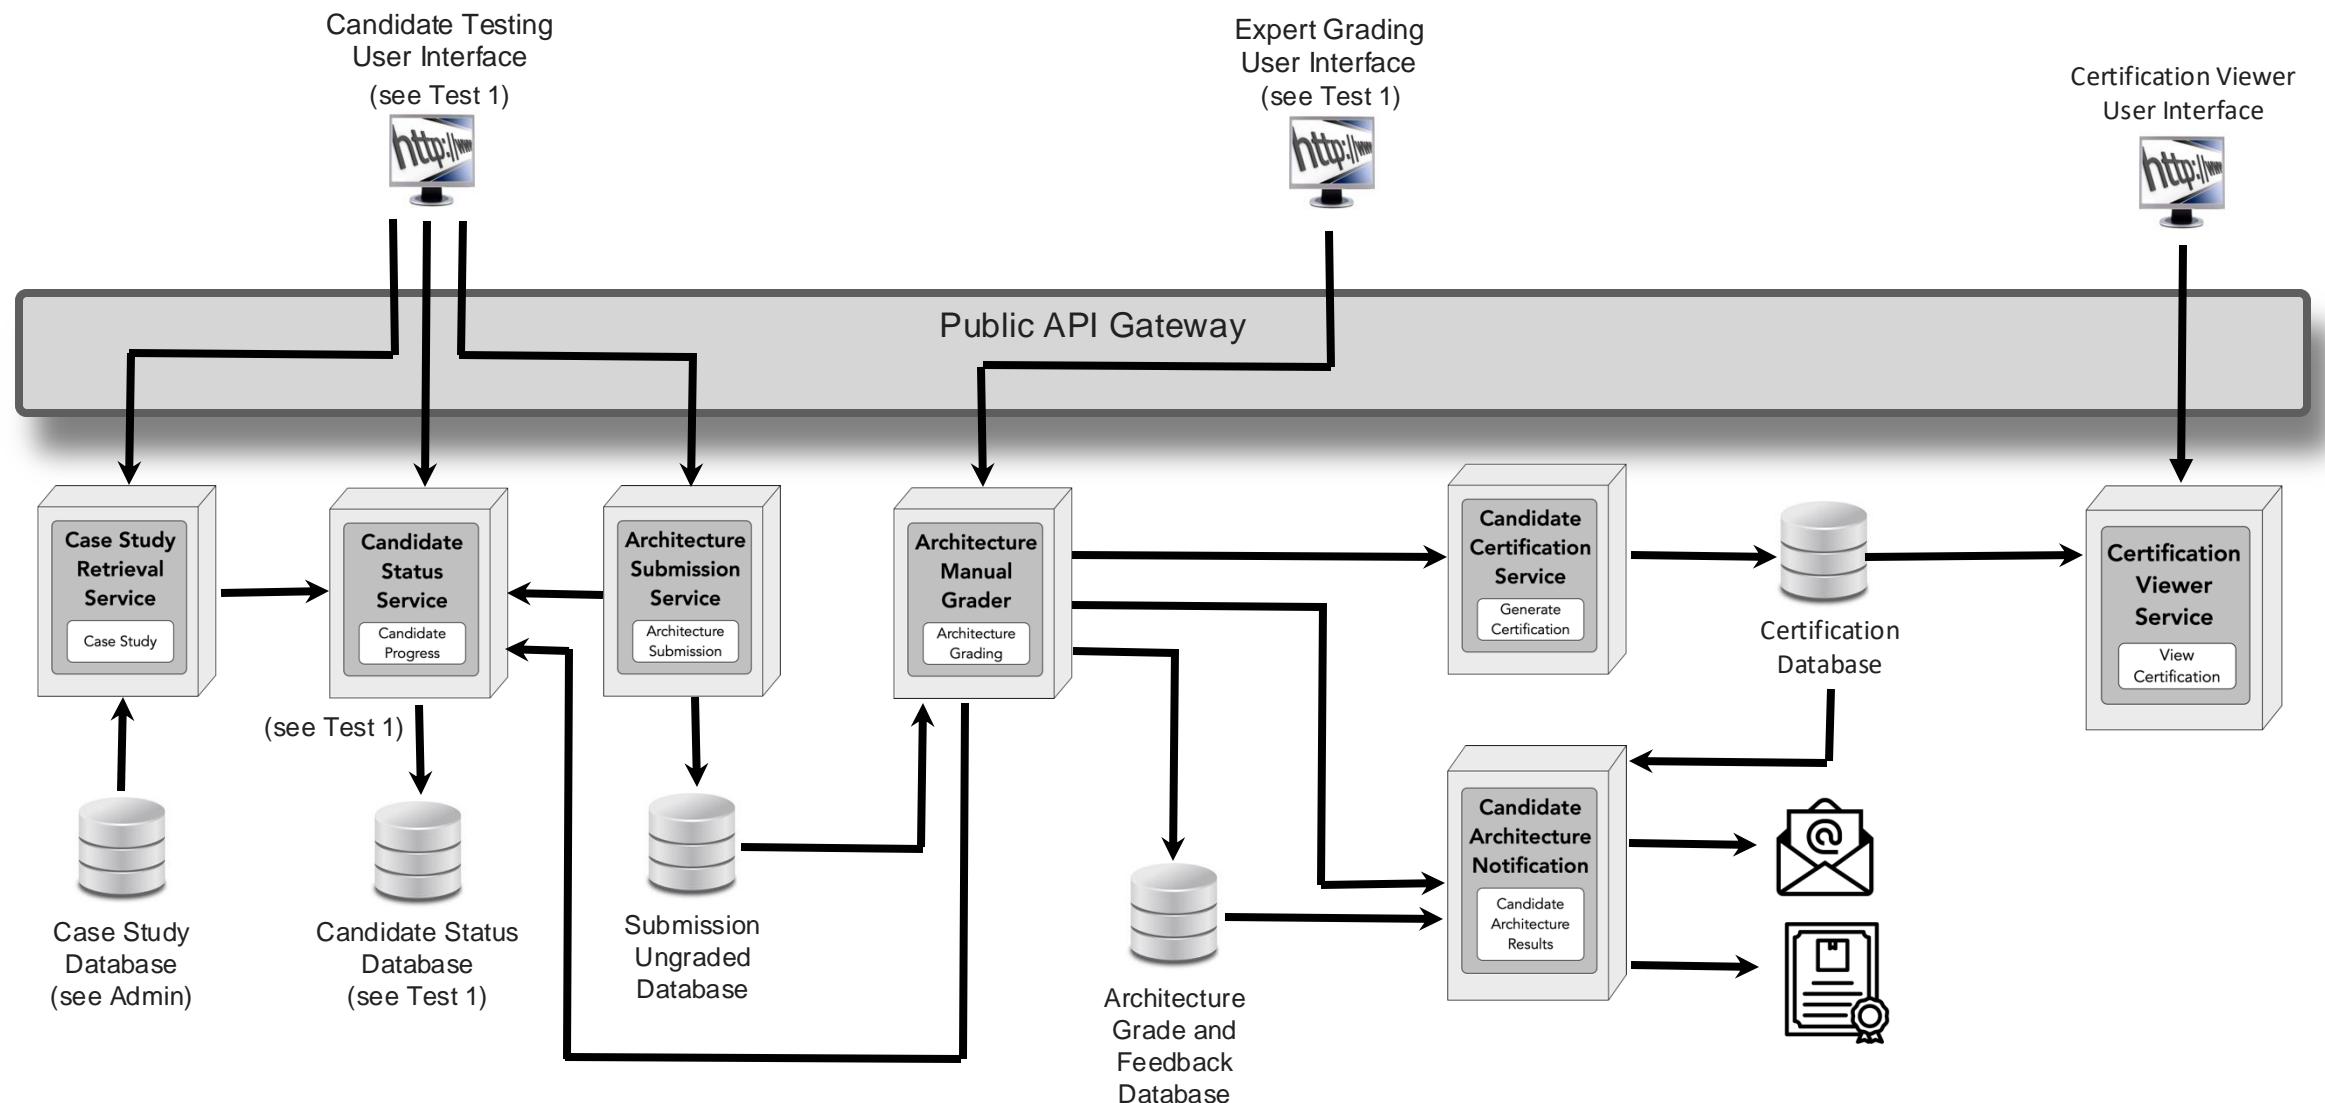
*Admin and expert sign-in artifacts removed for clarity

Existing Certification Testing Architecture (Test 1: Aptitude Test)



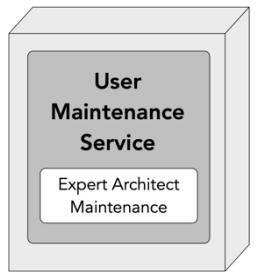
*Candidate, HR, and expert sign-in artifacts removed for clarity

Existing Certification Testing Architecture (Test 2: Architecture Solution)

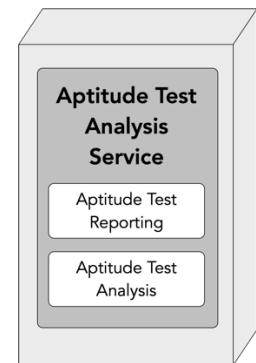


*Candidate, HR, and expert sign-in artifacts removed for clarity

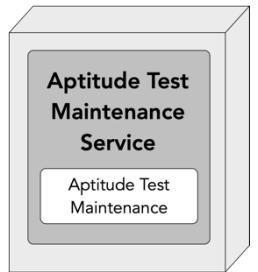
Existing Administrative Architecture



Add, remove, and maintain expert software architect consultants and their profiles, and used to identify “super experts” who are able to modify tests and case studies. Can be accessed by admin staff as well as designated experts.



Used to analyze aptitude certification test questions based on reports and analysis of question correctness percentages. Accessed by designated experts to periodically check the validity of the aptitude certification tests.



Add, remove, and modify aptitude test questions. Accessed by designated experts.

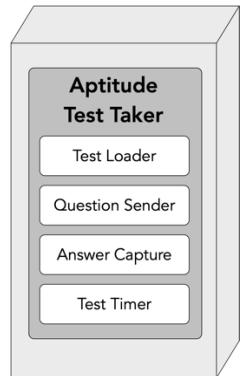


Add, remove, and modify architectural solution case studies. Accessed by designated experts.

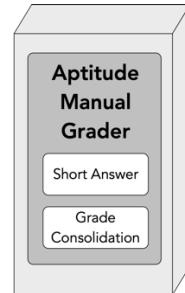
Existing Certification Testing Architecture (Test 1: Aptitude Test)



Used to register candidates for certification, and includes candidate profile information.

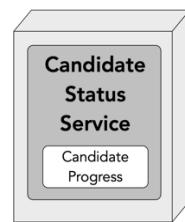


Primary service for administering aptitude tests. This service loads tests, delivers question to the user interface, captures candidate answers, and forwards these answers to the appropriate queue for processing as it receives them. Also times the test, and keeps track of where each candidate is in the certification test in the event of a restart or recovery situation.

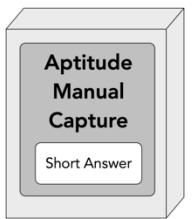


This service is used by the expert architects to retrieve ungraded short answer questions from the candidate and allows the expert to assign a grade as well as detailed feedback for each question.

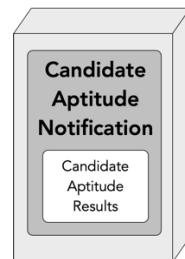
This service also acts as the main orchestrator to update the candidates certification status and also notify the candidate of the results.



Maintains the overall status of the candidates progress, including if they passed test 1 (aptitude), passed test 2 (architecture solution), and during test 2 the time that has elapsed since passing test 1 and retrieving the case study. Candidates have 30 days after passing test 1 to start test 2, and 2 weeks to complete test 2.

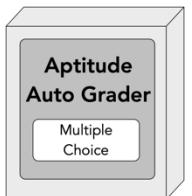


This service receives short answers from the queue and persists them in an ungraded answer database for later grading.



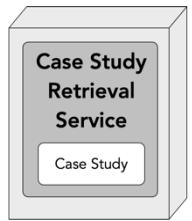
Notifies the candidate of their grade for test 1 (aptitude test) and also includes feedback on short answer questions from the expert if the answer was incorrect.

If the candidate passed test 1, the email also includes instructions for downloading the case study and starting test 2.



Receives multiple choice questions from a queue and automatically grades the answer based on the answer key in the test database.

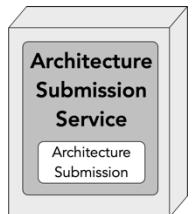
Existing Certification Testing Architecture (Test 2: Architecture Solution)



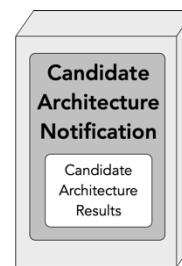
Used by the candidate to retrieve the assigned case study. Once the candidate downloads the case study test 2 officially starts, giving the candidate 2 weeks to upload their solution. If 30 days has passed since being notified that they passed test 1, the candidate will be unable to download the case study and will have to start over.



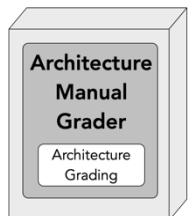
This service generates the official certification certificate and inserts the certification into the certification database.



Used by the candidate to upload their architectural solution.



Notifies the candidate of their grade for the architecture submission and also includes feedback from the expert justifying the grade. It also notifies the candidate whether they passed the certification or not, and includes further instructions.



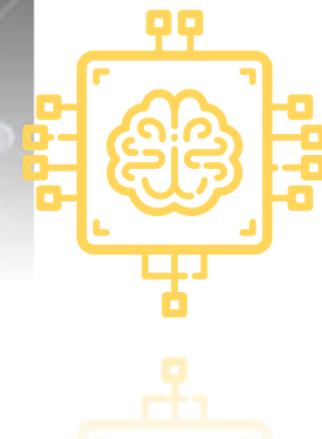
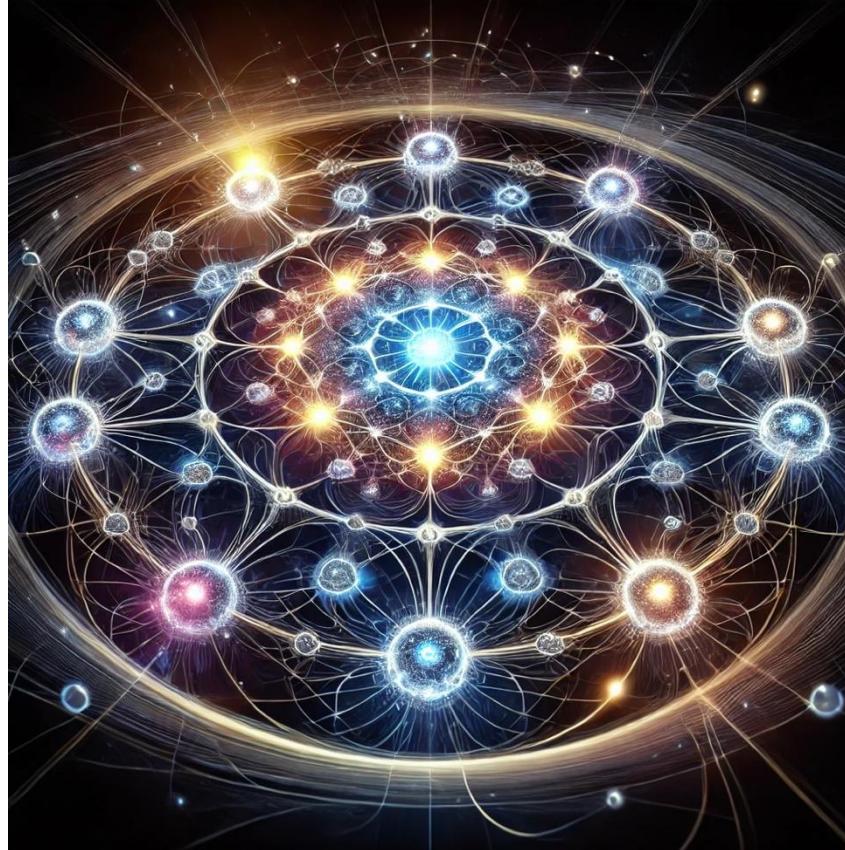
Used by the experts to review the architecture solution from the candidate and grade the solution. The expert also provides feedback to the candidate using this service.

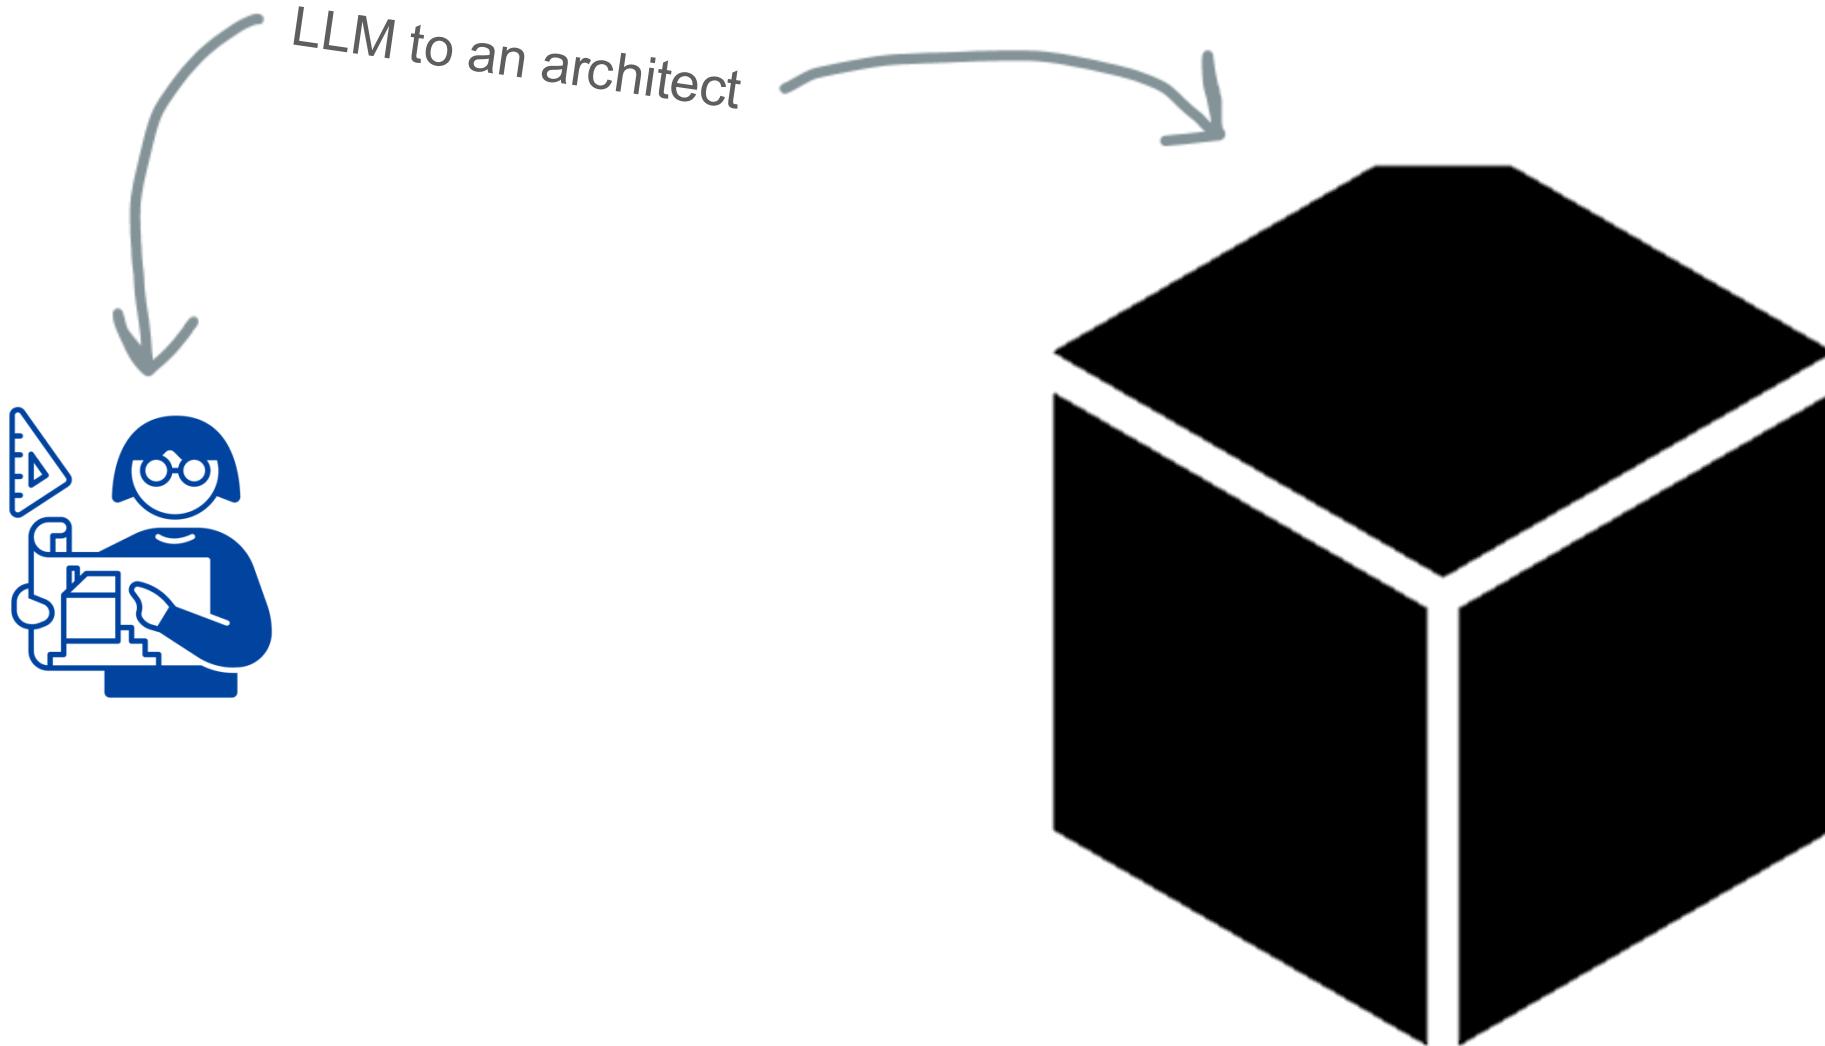


This service allows the candidate and company HR representatives to view and verify the certification status for an individual.

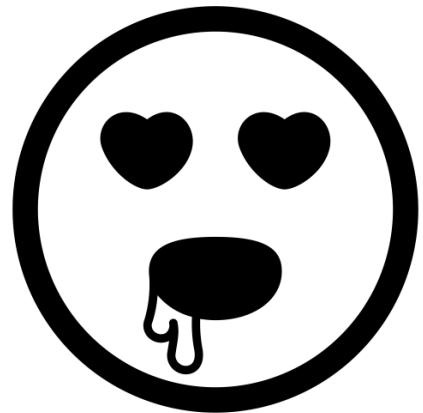
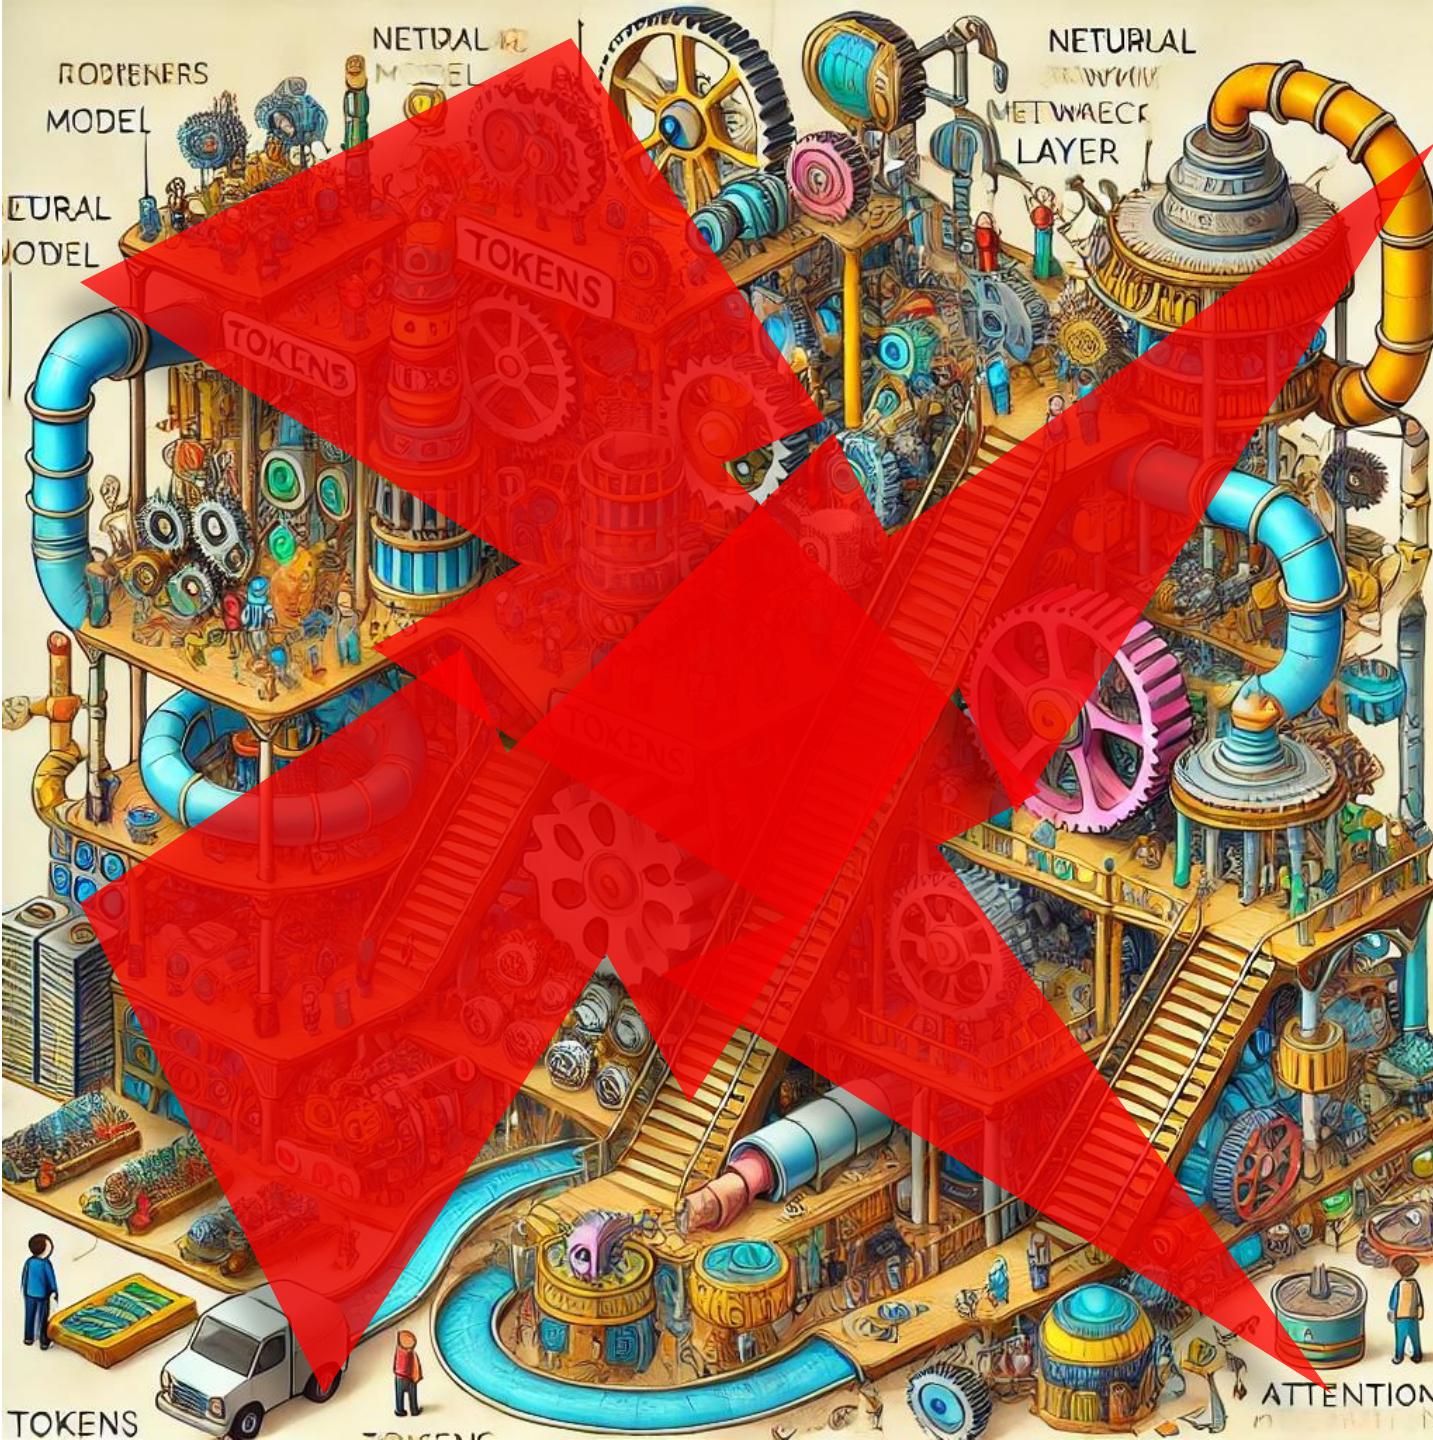
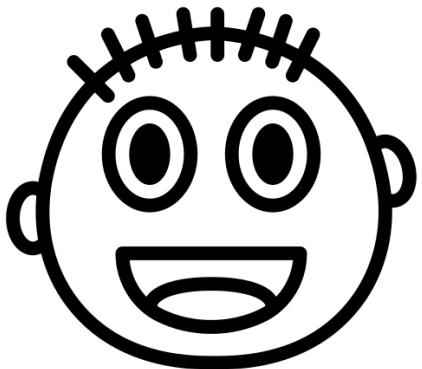
This service also acts as the main orchestrator to update the candidate's certification status and also notify the candidate of the results, and if they pass to generate the certification.

Generative AI Architecture





Of course, generated
by AI

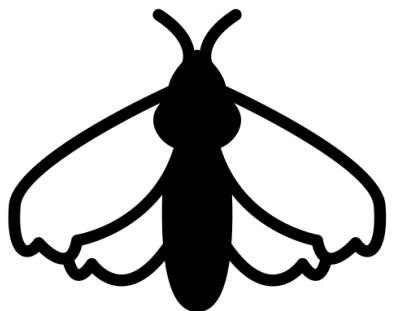


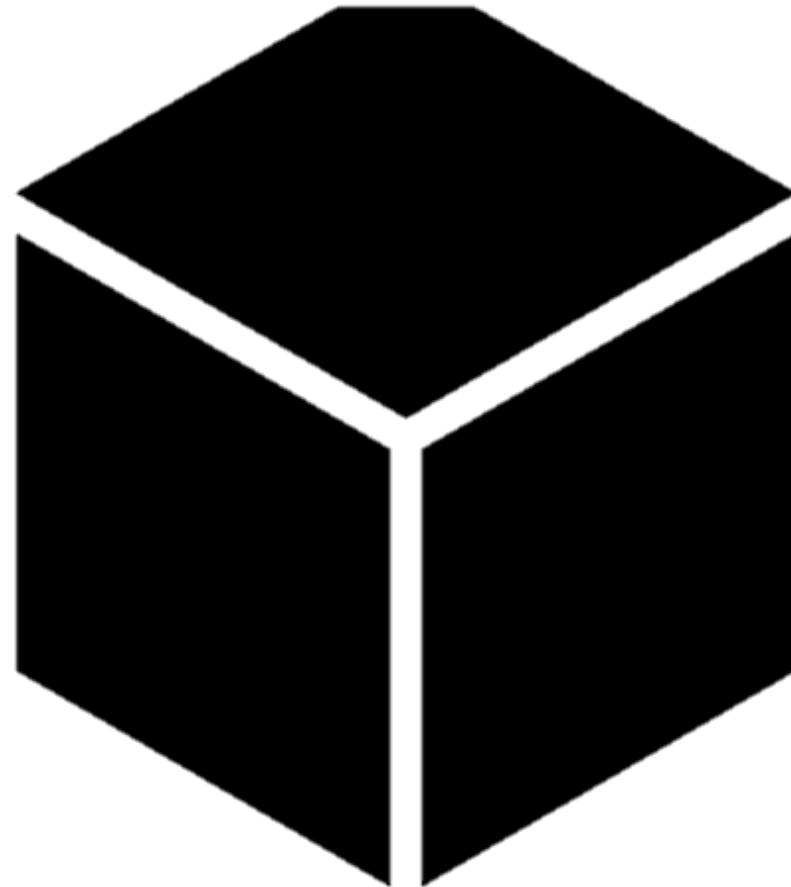


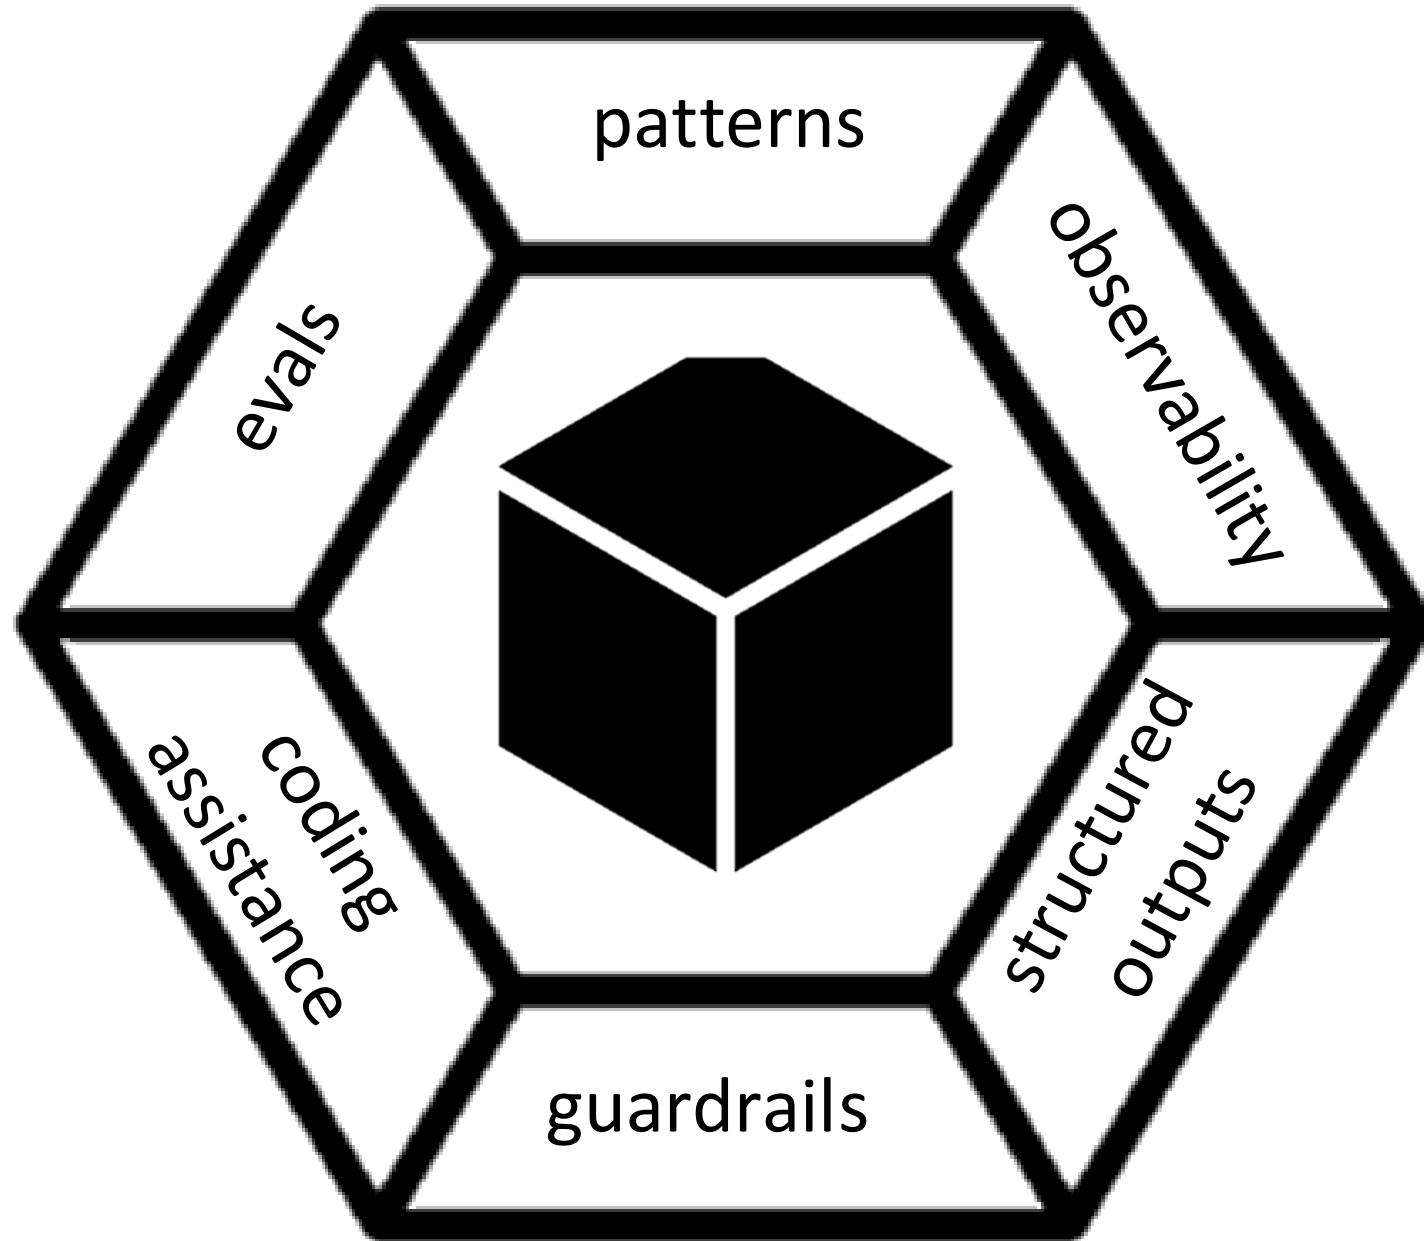
architects, too

“Developers are drawn to complexity like moths to a flame, frequently with the same results.”

Neal, writing in *97 Things Every Architect Should Know*







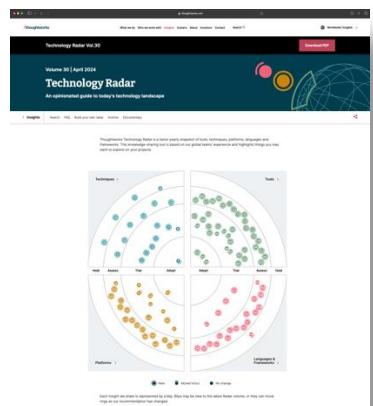
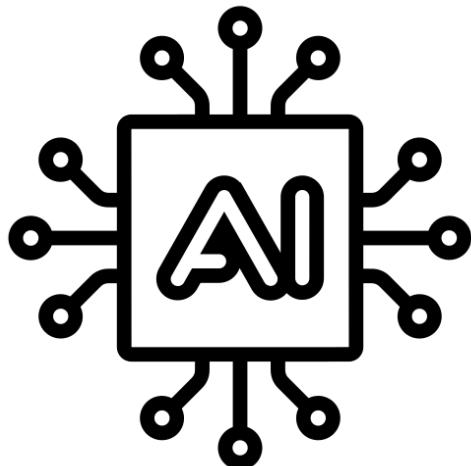
Pre final cull: 51

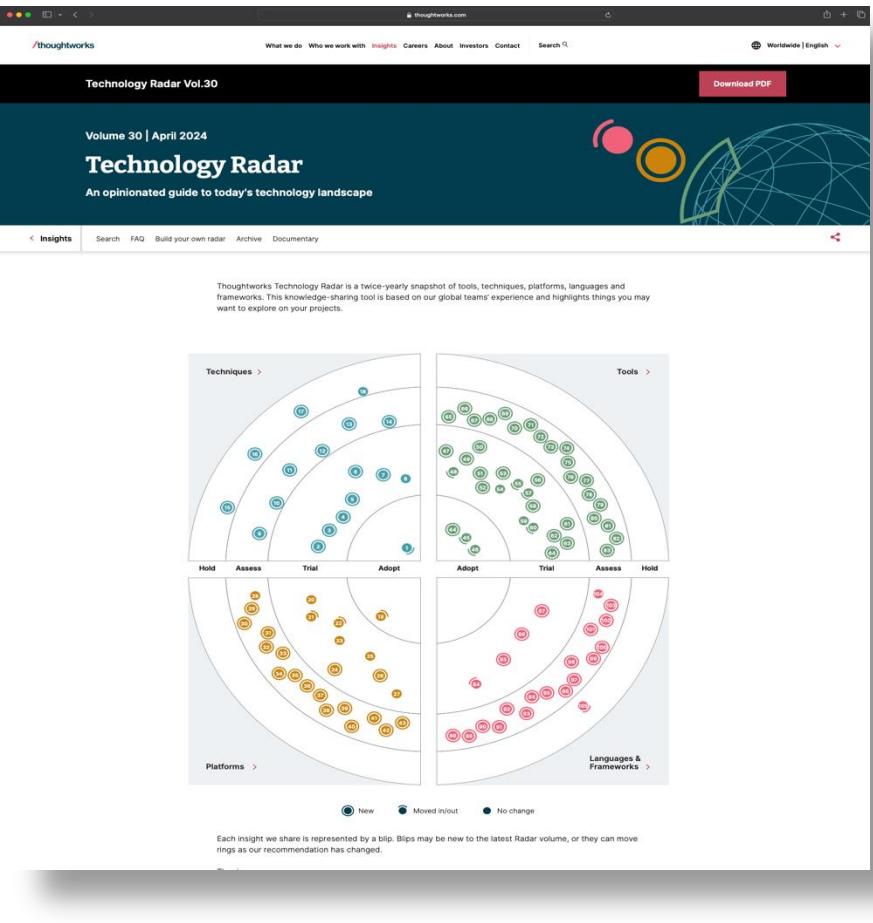
Post final cull: 42

AI-assistance 8, rest 36

Post life boat: 43 (38%)

AI in Architecture


www.thoughtworks.com/radar

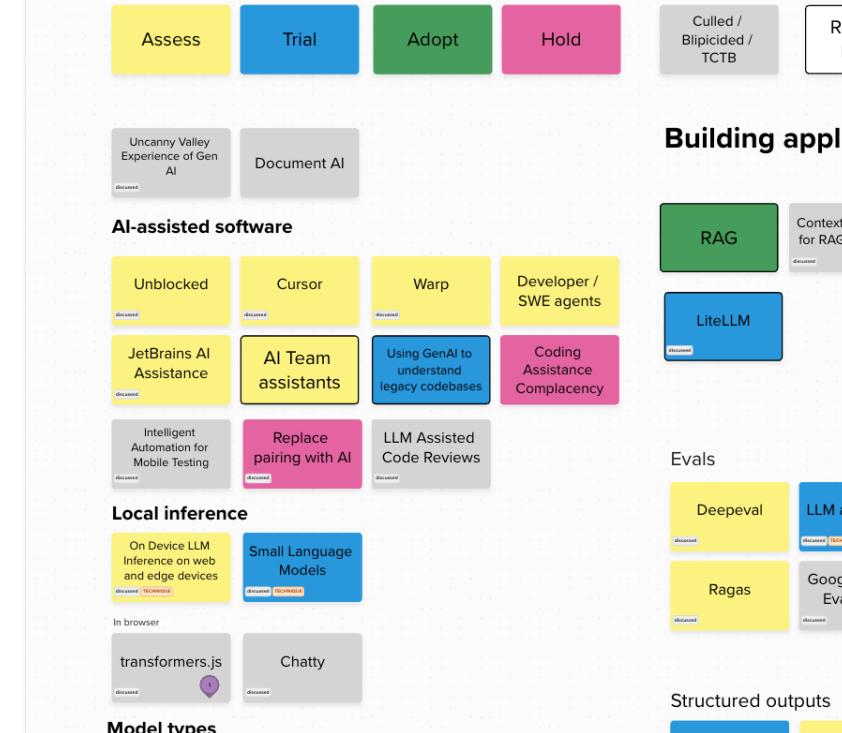
www.thoughtworks.com/radar

vol31

Pre final cull: 51

Post final cull: 42

AI-assistance 8, rest 36



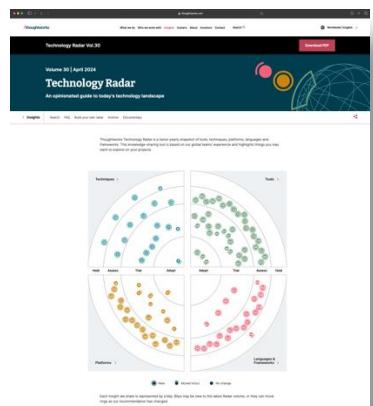
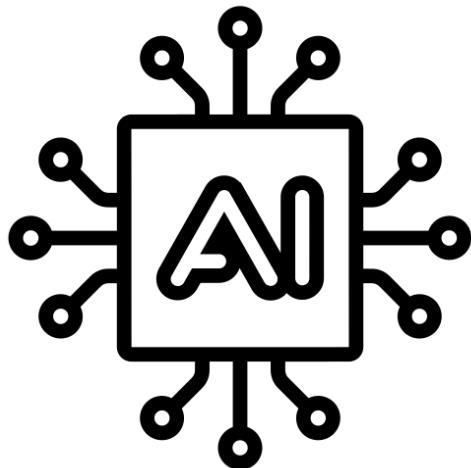
Pre final cull: 51

Post final cull: 42

AI-assistance 8, rest 36

Post life boat: 43 (38%)

AI in Architecture



www.thoughtworks.com/radar



vol31

Pre final cull: 51

Post final cull: 42

AI-assistance 8, rest 36

Post life boat: 43 (38%)

Assess

Trial

Adopt

Hold

Culled /
Blipicided /
TCTB

Reblip /
Move

Uncanny Valley
Experience of Gen
AI

Document AI

AI-assisted software

Unblocked

Cursor

Warp

Developer /
SWE agents

JetBrains AI
Assistance

AI Team
assistants

Using GenAI to
understand
legacy codebases

Coding
Assistance
Complacency

Intelligent
Automation for
Mobile Testing

Replace
pairing with AI

LLM Assisted
Code Reviews

Local inference

Building applications with LLMs

RAG

Context Caching
for RAG (Gemini)

kotaemon

LiteLLM

LlamaIndex

Function calling
with LLMs

Gorilla

Semantic
Router

Evals

Deepeval

LLM as a judge

Guardrails

NeMo

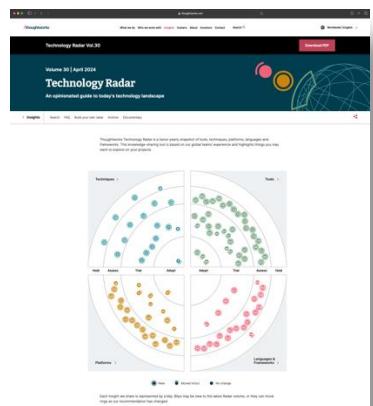
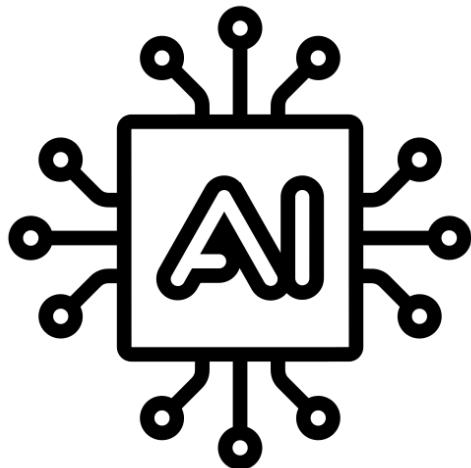
Pre final cull: 51

Post final cull: 42

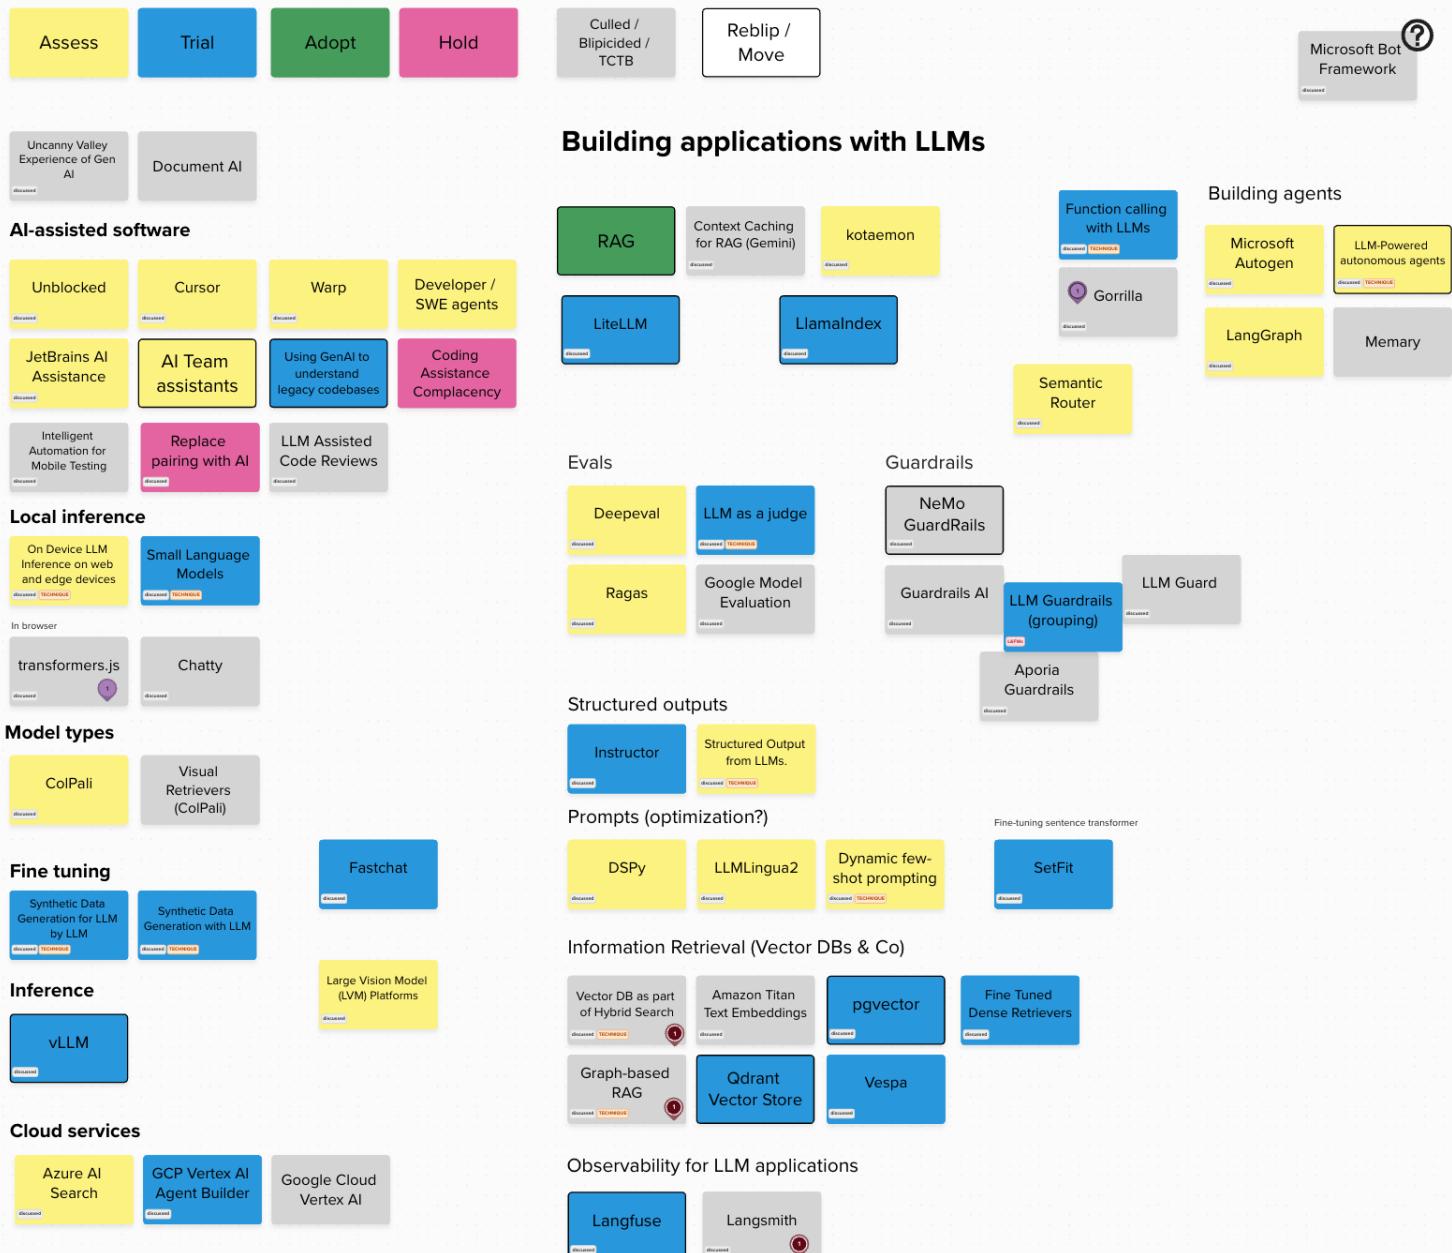
AI-assistance 8, rest 36

Post life boat: 43 (38%)

AI in Architecture



www.thoughtworks.com/radar



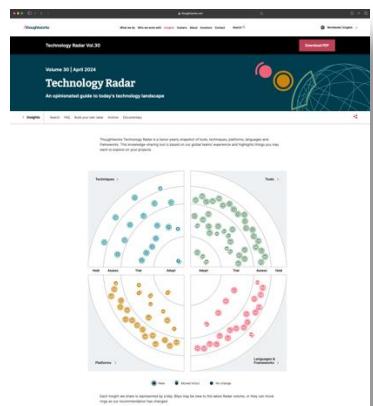
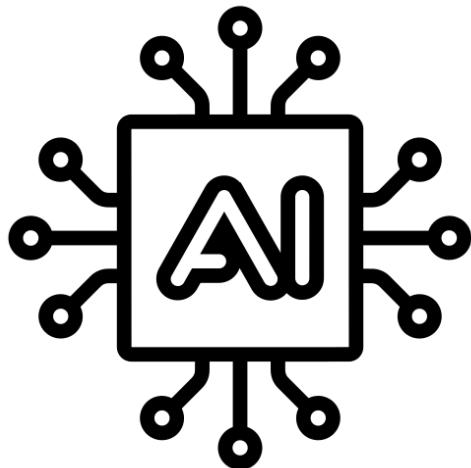
Pre final cull: 51

Post final cull: 42

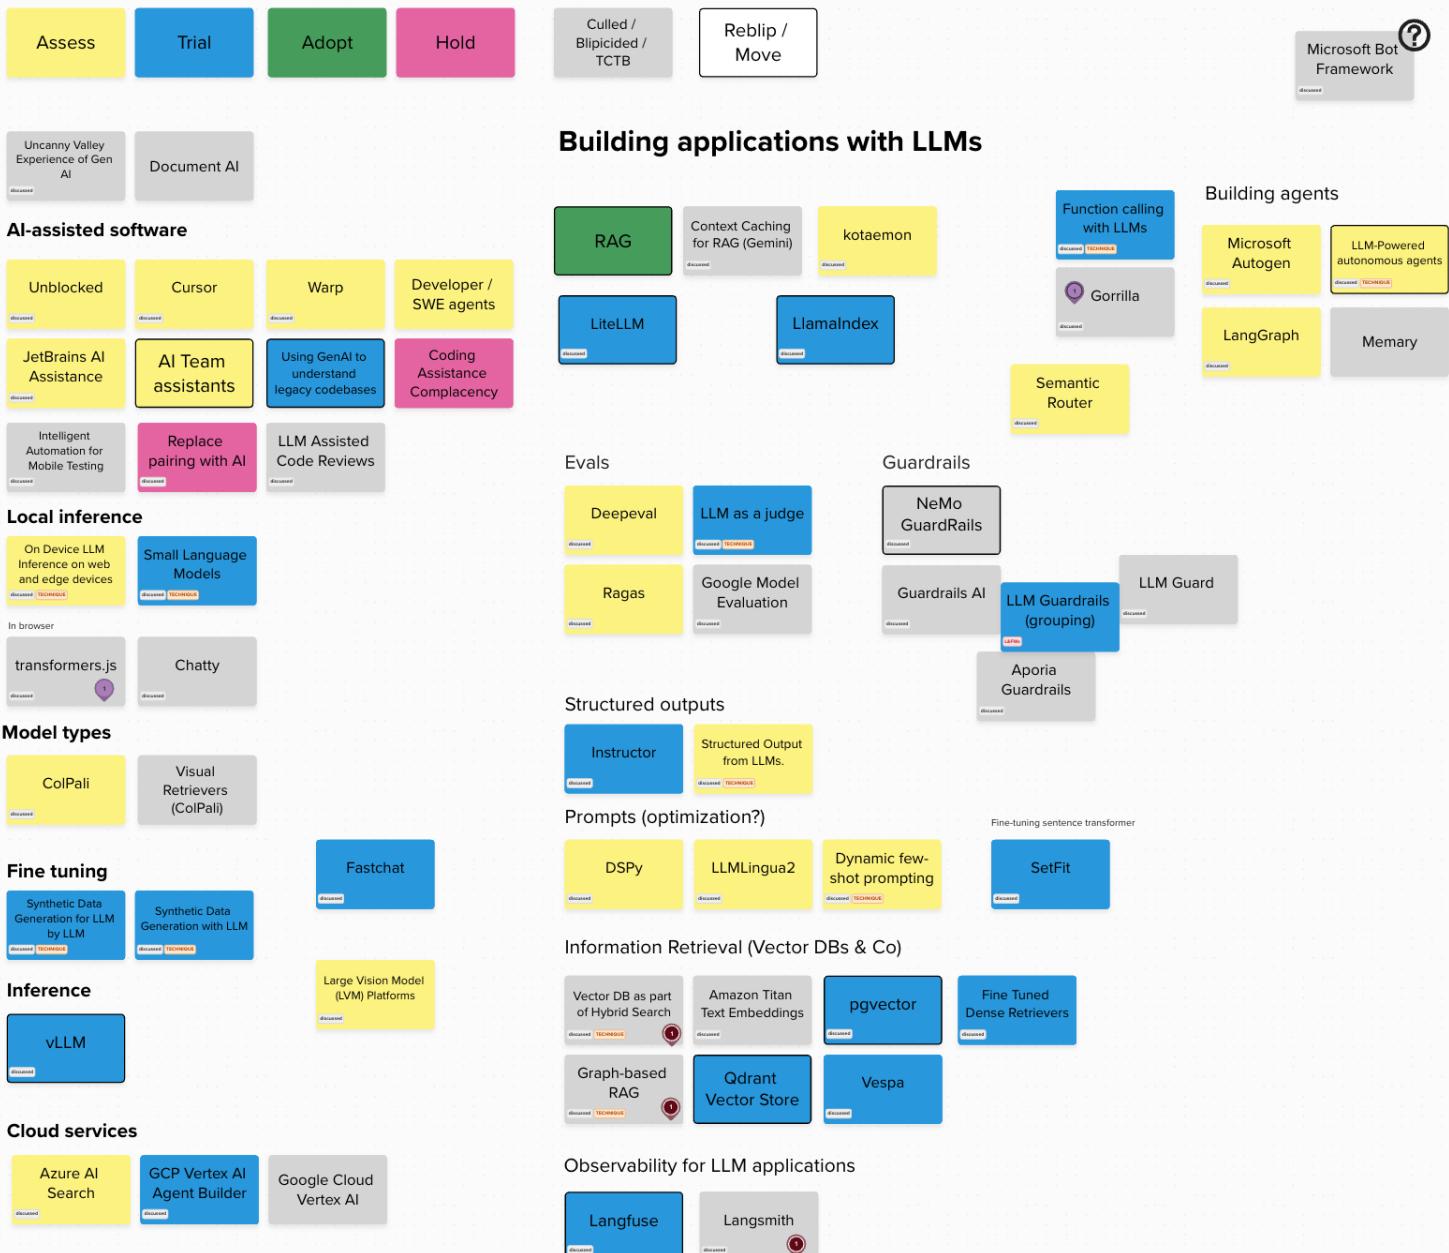
AI-assistance 8, rest 36

Post life boat: 43 (38%)

AI in Architecture



www.thoughtworks.com/radar



Pre final cull: 51

Post final cull: 42

AI-assistance 8, rest 36

Post life boat: 43 (38%)



Categories:

AI-assisted software dev

Local inference

Fine Tuning

Inference

Cloud services

Evals/Guardrails

Structured outputs

Prompts

Information retrieval

Observability for LLM

Building agents

Pre final cull: 51

Post final cull: 42

AI-assistance 8, rest 36

Post life boat: 43 (38%)



Categories:

AI-assisted software dev

Local inference

Fine Tuning

Inference

Cloud services

Evals/Guardrails

Structured outputs

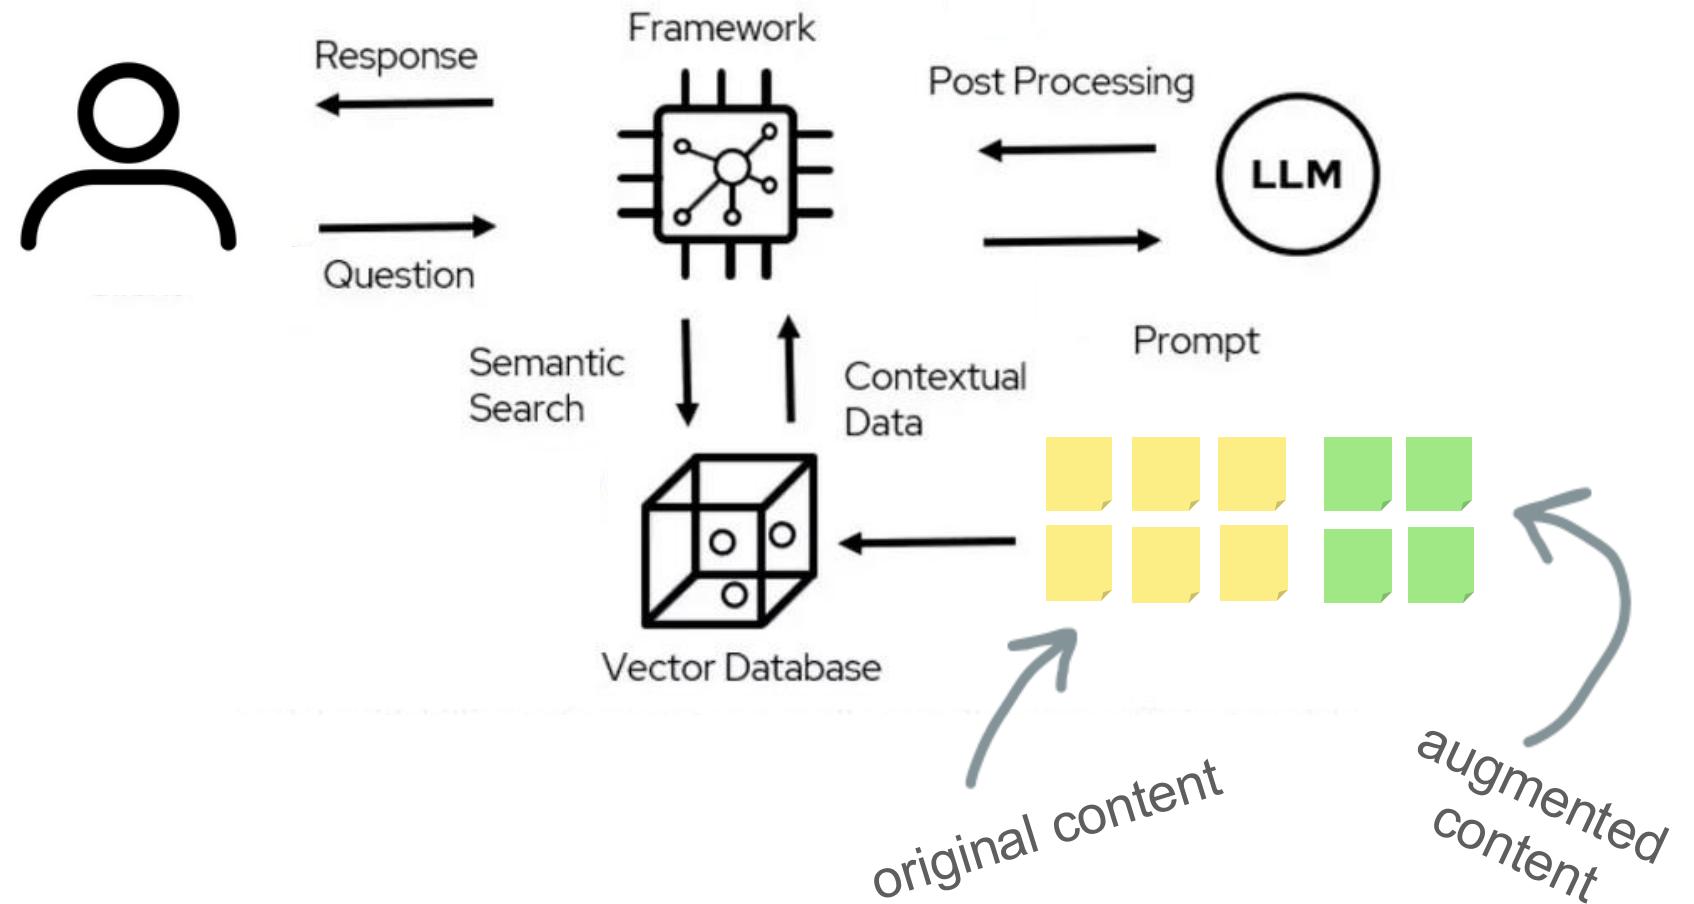
Prompts

Information retrieval

Observability for LLM

Building agents

RAG



Pre final cull: 51

Post final cull: 42

AI-assistance 8, rest 36

Post life boat: 43 (38%)



Categories:

AI-assisted software dev

Local inference

Fine Tuning

Inference

Cloud services

Evals/Guardrails

Structured outputs

Prompts

Information retrieval

Observability for LLM

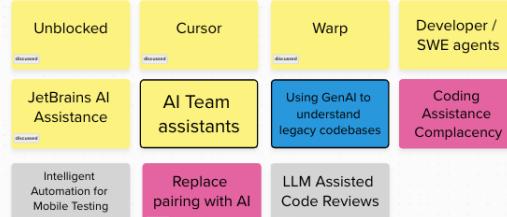
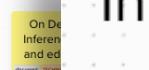
Building agents

Pre final cull: 51

Post final cull: 42

AI-assistance 8, rest 36

Post life boat: 43 (38%)

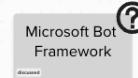
**AI-assisted software****Local****Model****Fine tune****Inference****Cloud services****Building applications with LLMs**

Evals

Guardrails



Building agents

**Information Retrieval (Vector DBs & Co)**

Vector DB as part of Hybrid Search

Amazon Titan Text Embeddings

pgvector

Fine Tuned Dense Retrievers

Graph-based RAG

Qdrant Vector Store

Vespa

Observability for LLM applications**Categories:**

AI-assisted software dev

Local inference

Fine Tuning

Inference

Cloud services

Evals/Guardrails

Structured outputs

Prompts

Information retrieval

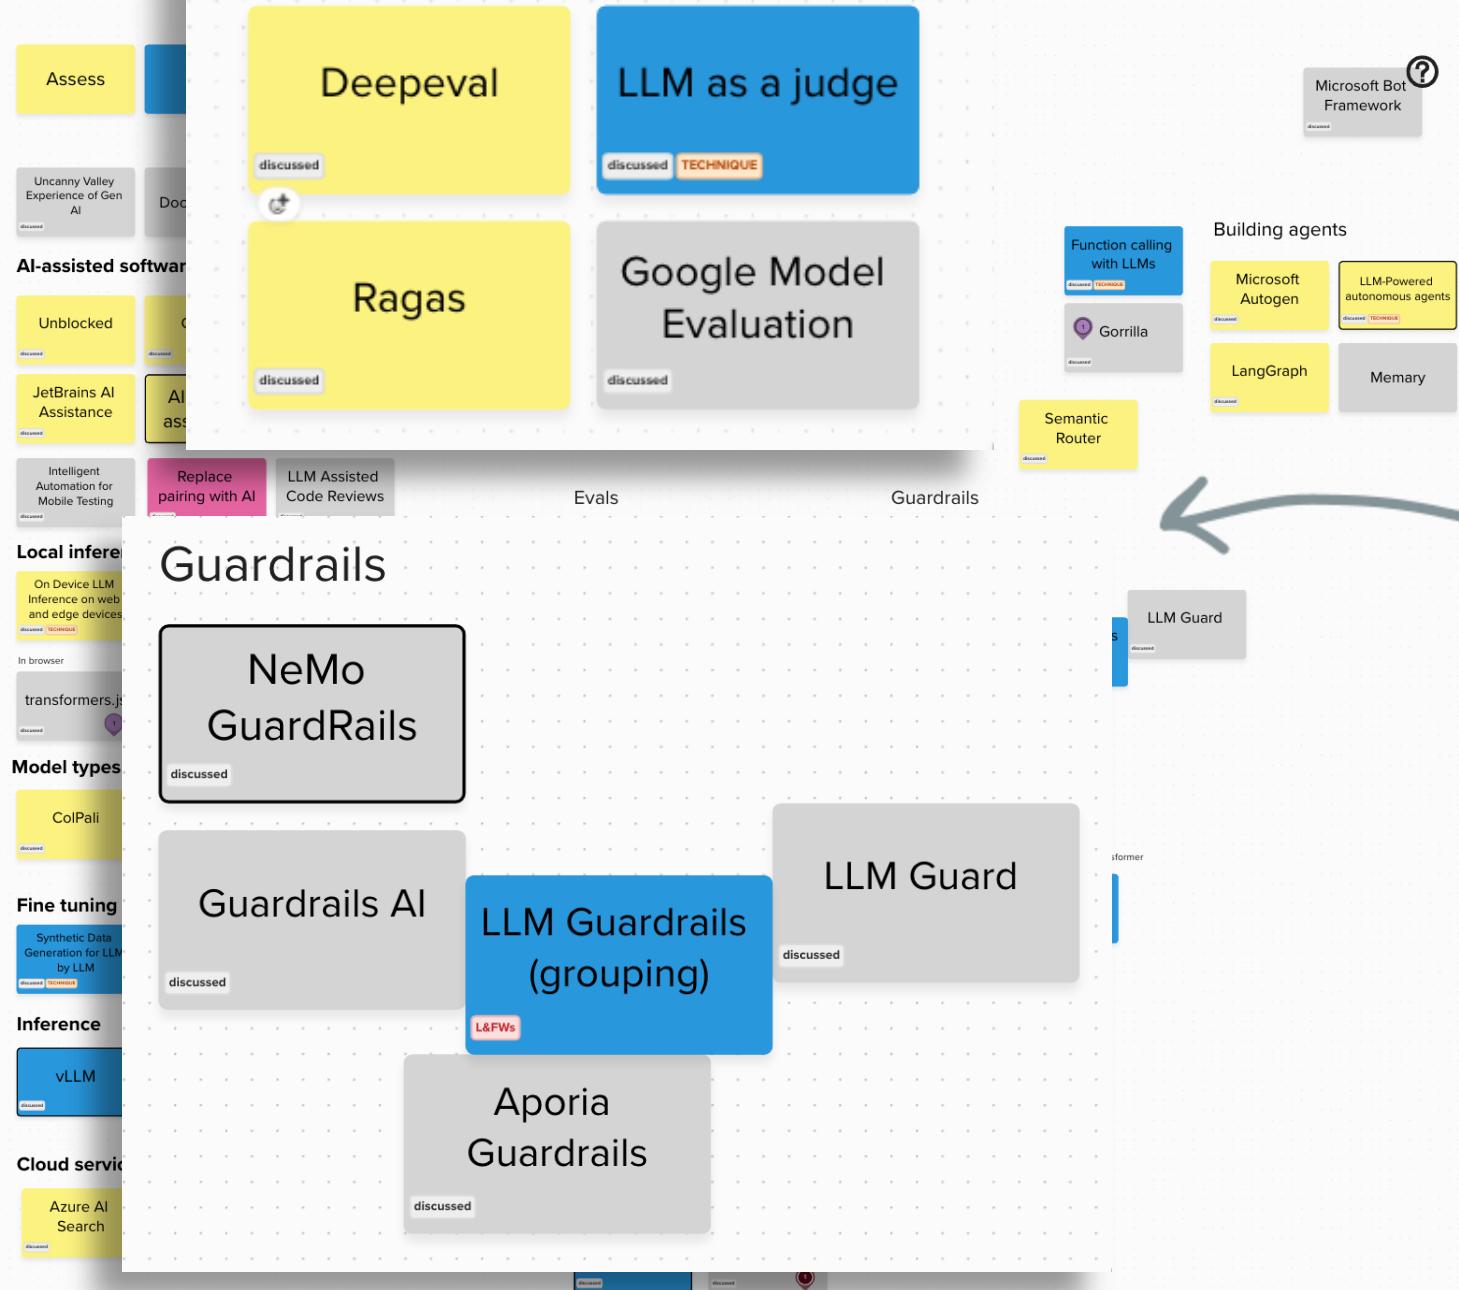
Observability for LLM

Building agents

Pre final c

Evals

boat: 43 (38%)



Categories:

AI-assisted software dev

Local inference

Fine Tuning

Inference

Cloud services

Evals/Guardrails

Structured outputs

Prompts

Information retrieval

Observability for LLM

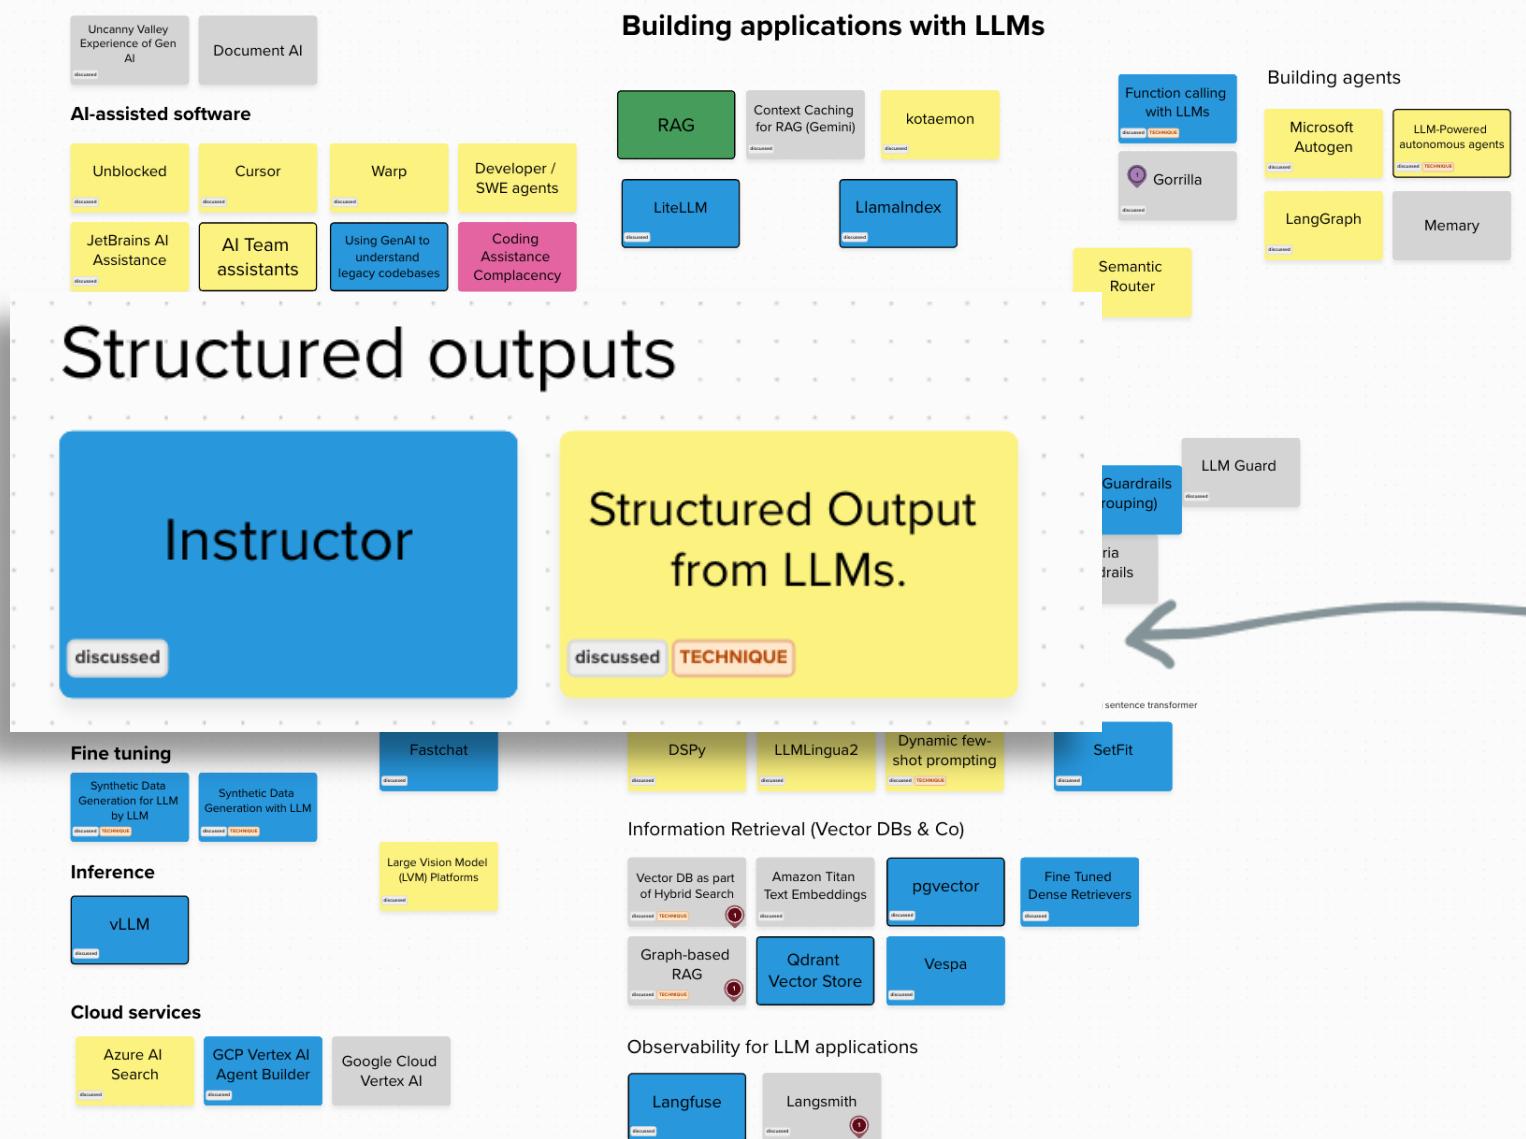
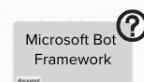
Building agents

Pre final cull: 51

Post final cull: 42

AI-assistance 8, rest 36

Post life boat: 43 (38%)

**Categories:**

AI-assisted software dev

Local inference

Fine Tuning

Inference

Cloud services

Evals/Guardrails

Structured outputs

Prompts

Information retrieval

Observability for LLM

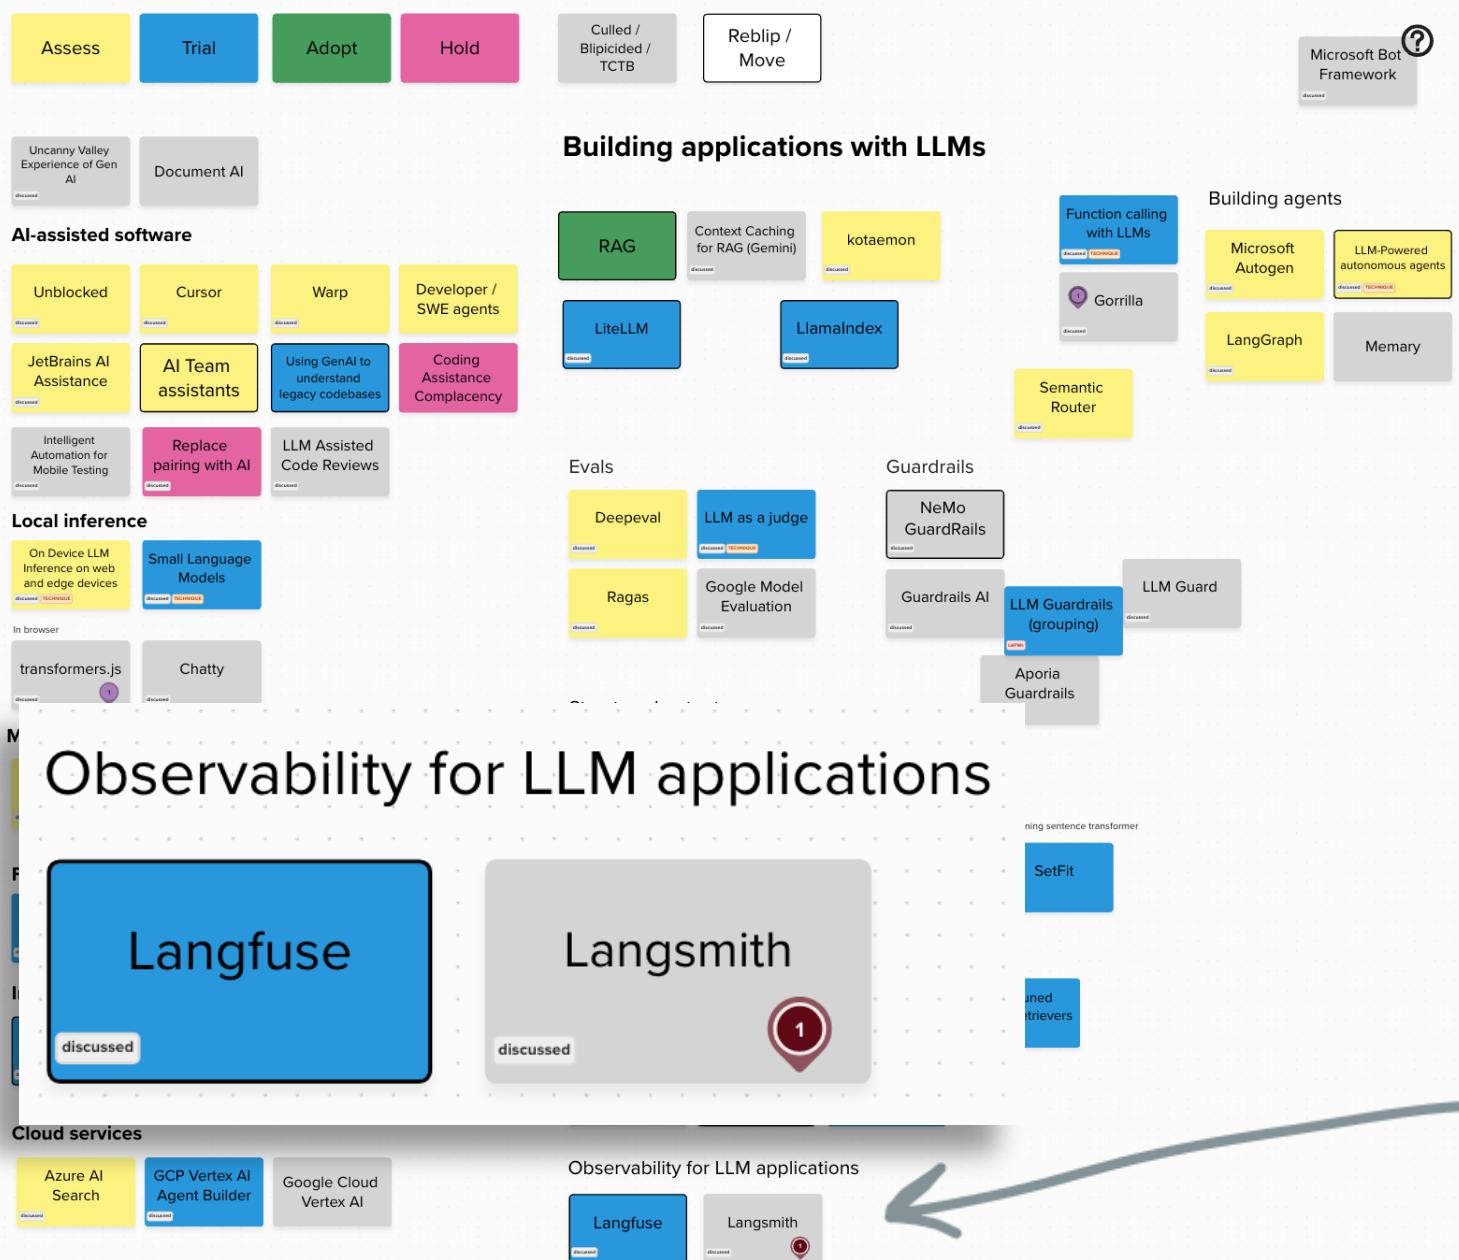
Building agents

Pre final cull: 51

Post final cull: 42

AI-assistance 8, rest 36

Post life boat: 43 (38%)



Categories:

AI-assisted software dev

Local inference

Fine Tuning

Inference

Cloud services

Evals/Guardrails

Structured outputs

Prompts

Information retrieval

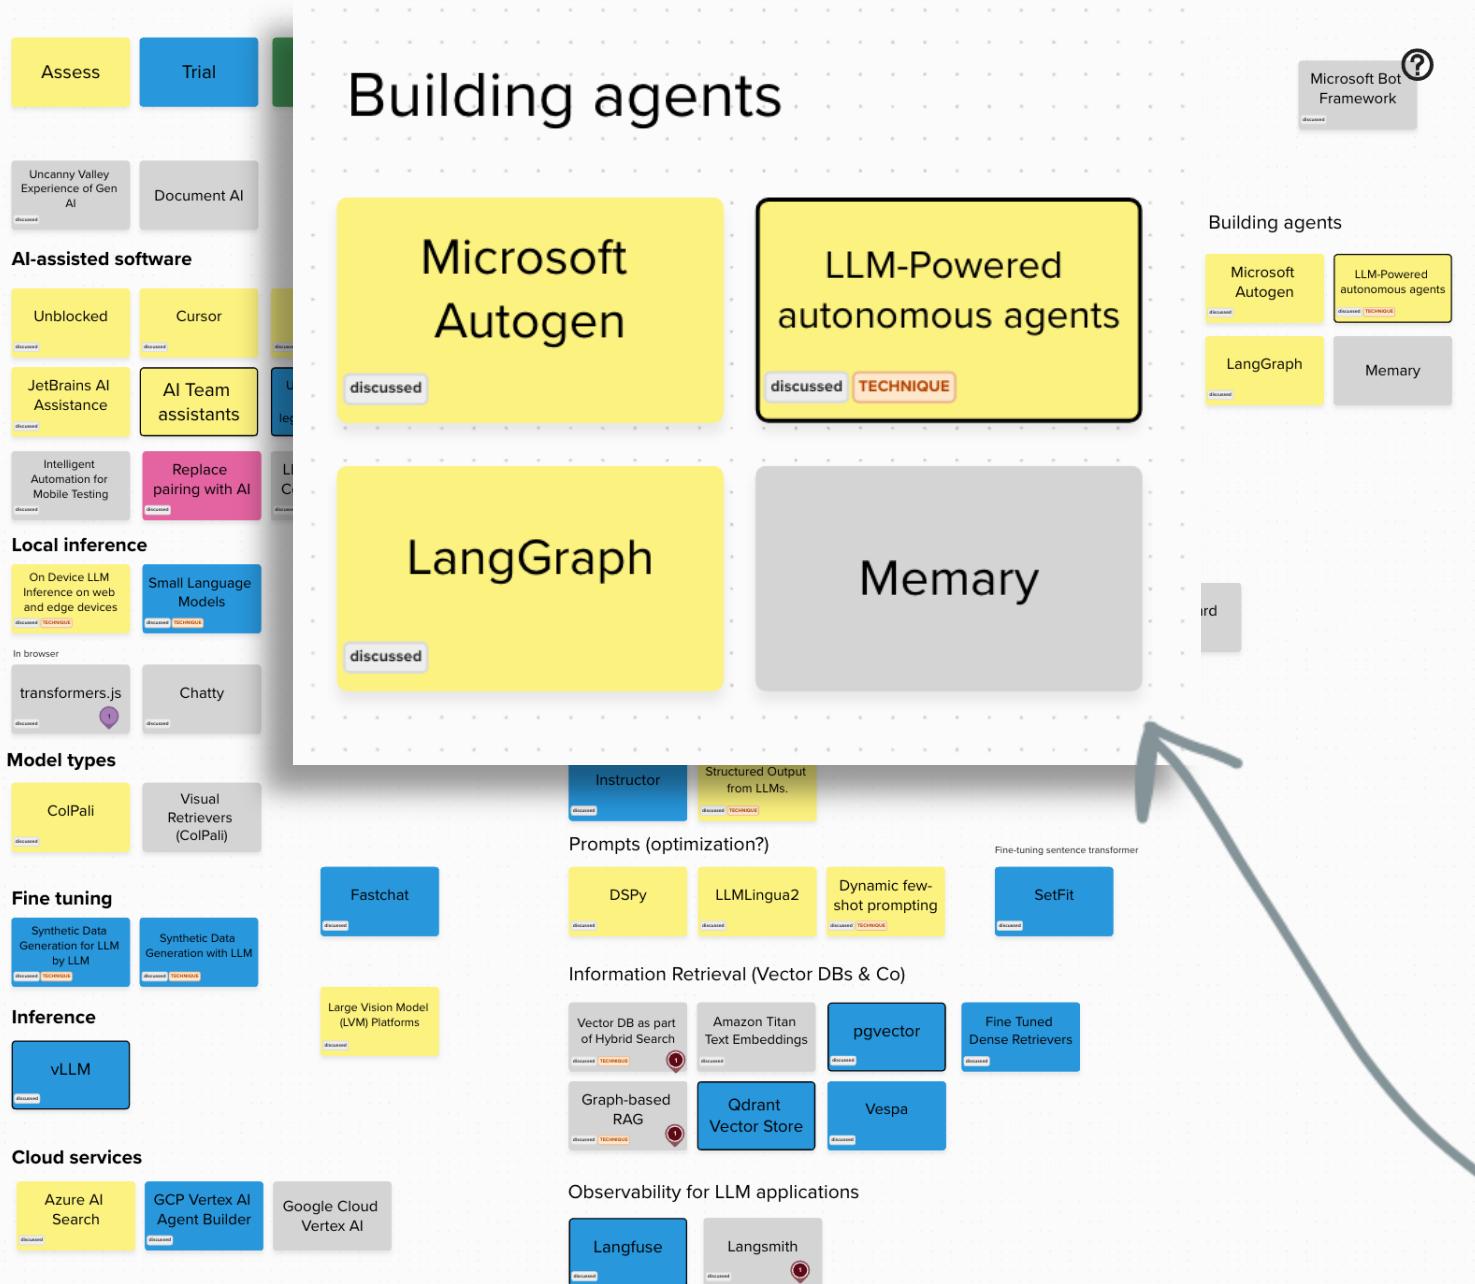
Observability for LLM

Building agents

Pre final cull: 51

Post final cull: 42
AI-assistance 8, rest 36

Post life boat: 43 (38%)



Categories:

AI-assisted software dev

Local inference

Fine Tuning

Inference

Cloud services

Evals/Guardrails

Structured outputs

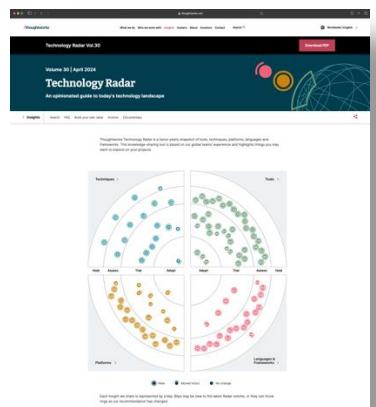
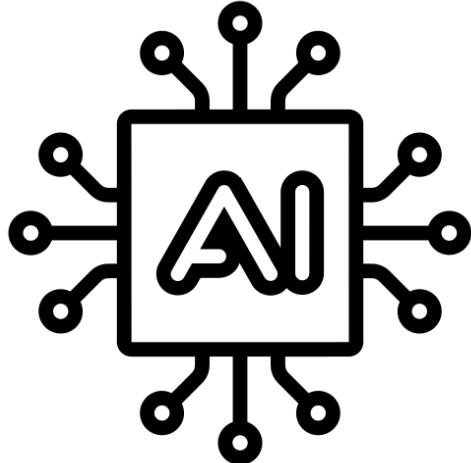
Prompts

Information retrieval

Observability for LLM

Building agents

AI n Architecture



www.thoughtworks.com/radar

vol31

Pre final cull: 51

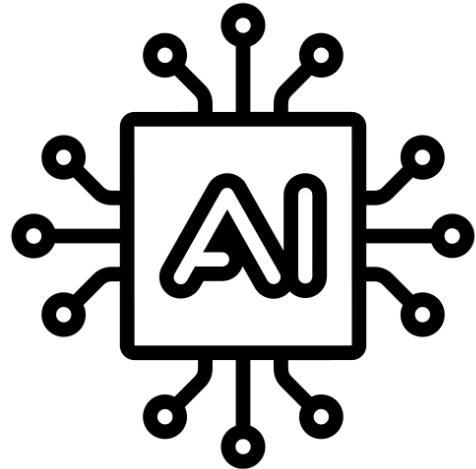
Post final cull: 42

AI-assistance 8, rest 36

Post life boat: 43 (38%)



AI ∩ Architecture



Components in LLM-backed architectures

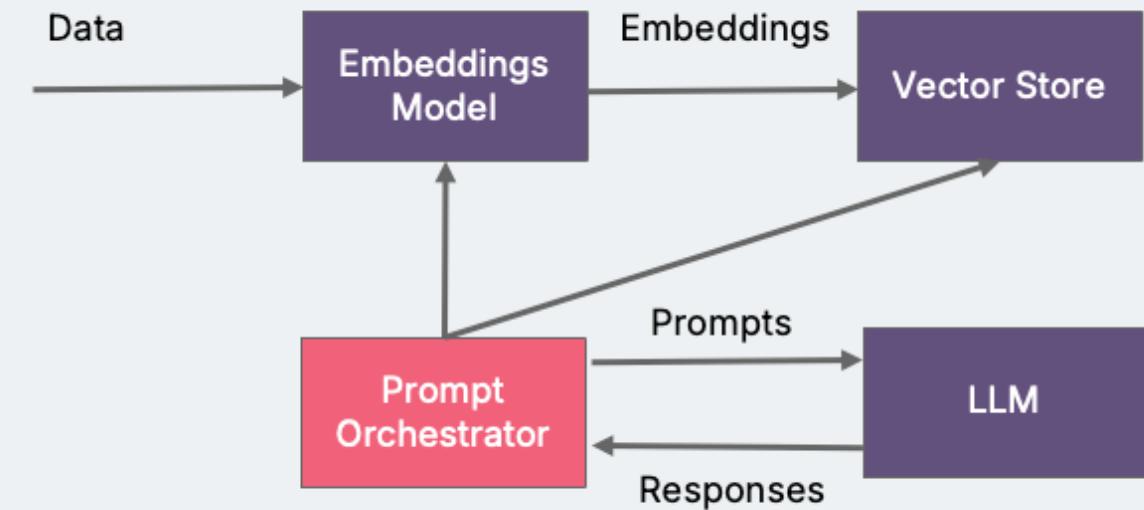


**Let's send
prompts to an
LLM**

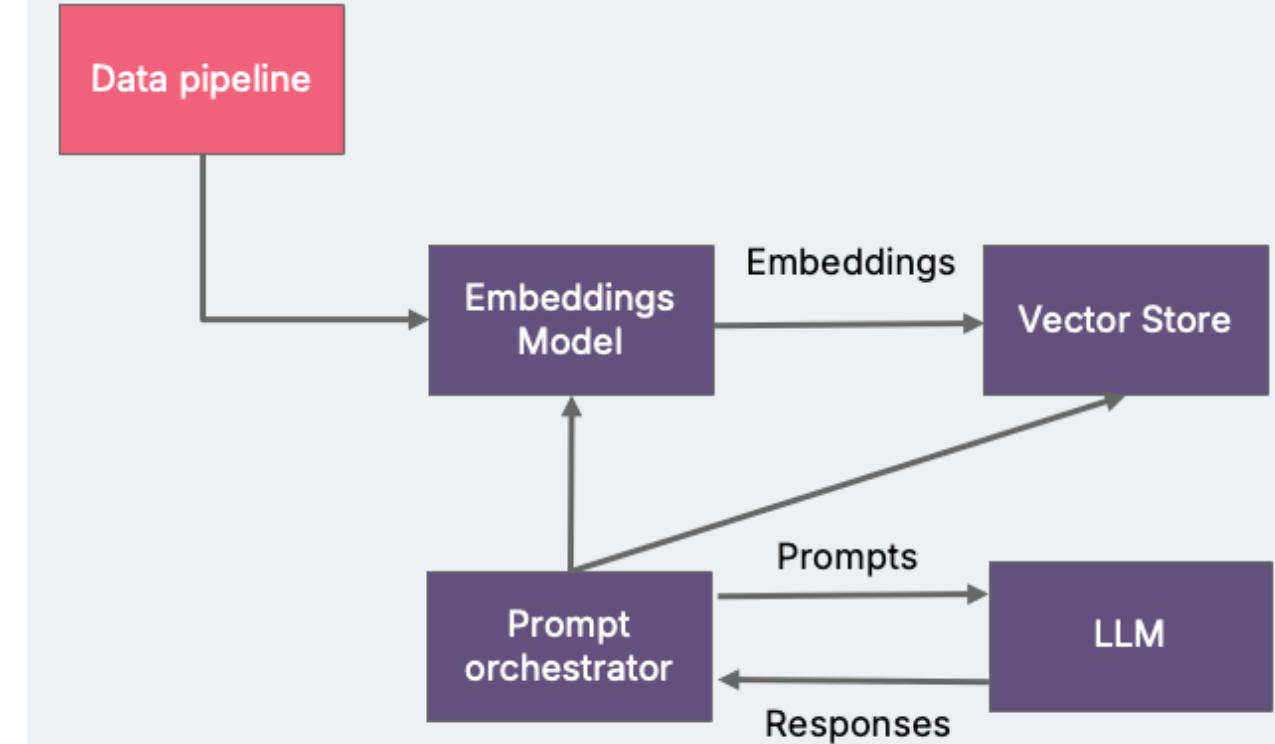


**We want the
LLM to know
about our data**

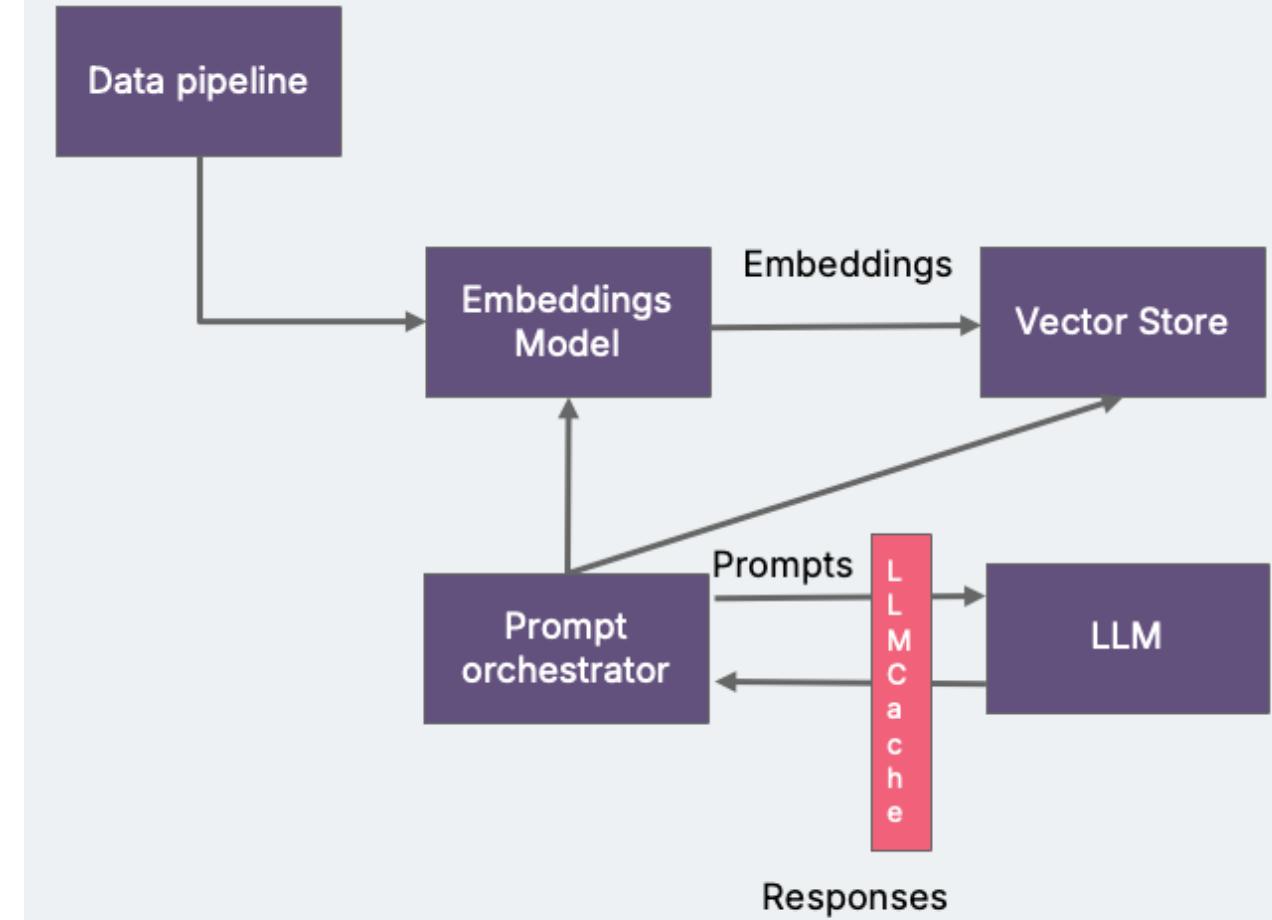
**#retrieval
Augmented
Generation**



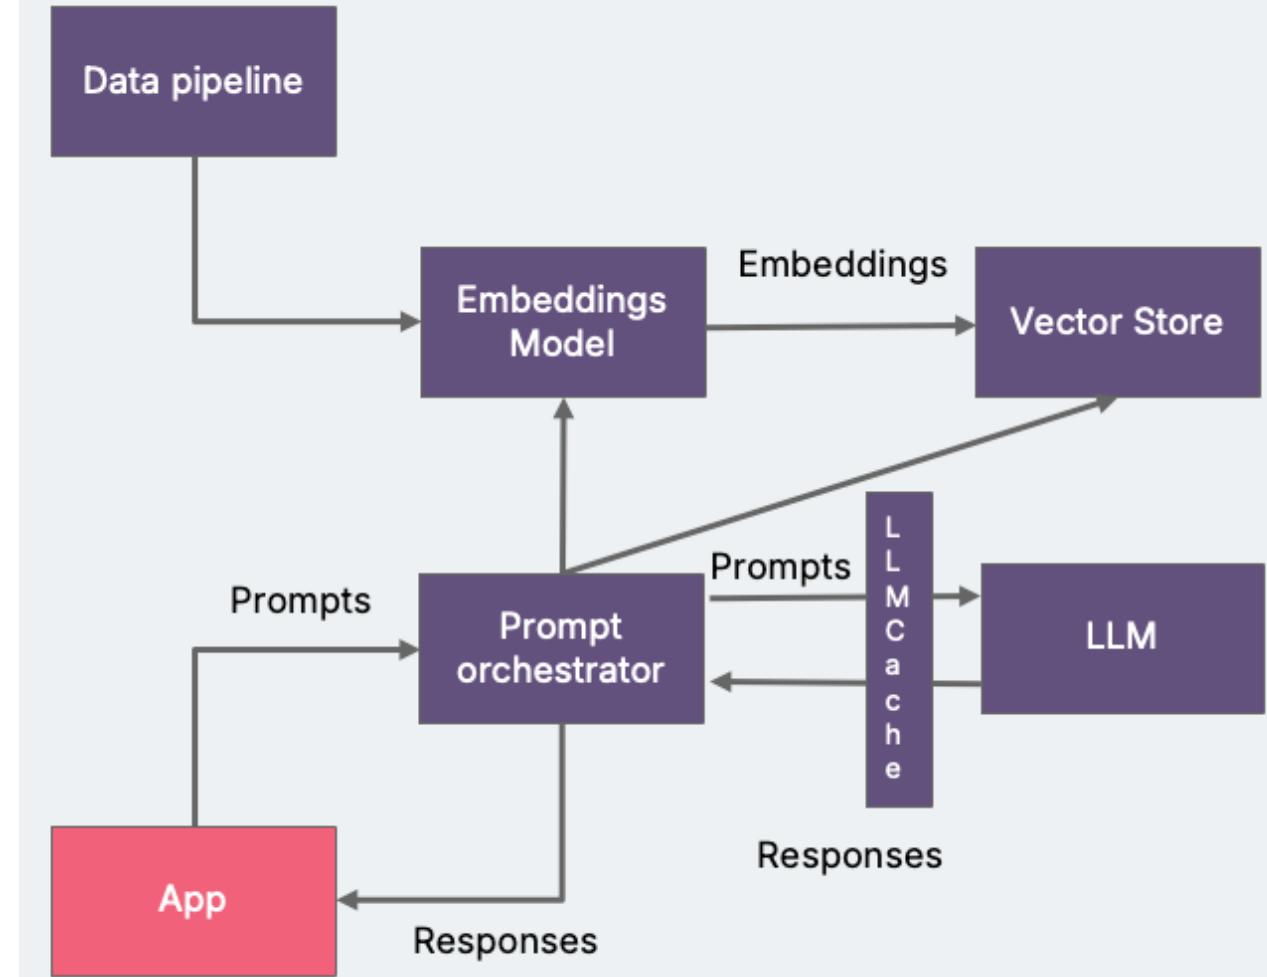
**We want to have
up-to-date data
available**



**Do we want to
hit the LLM all
the time
though?**

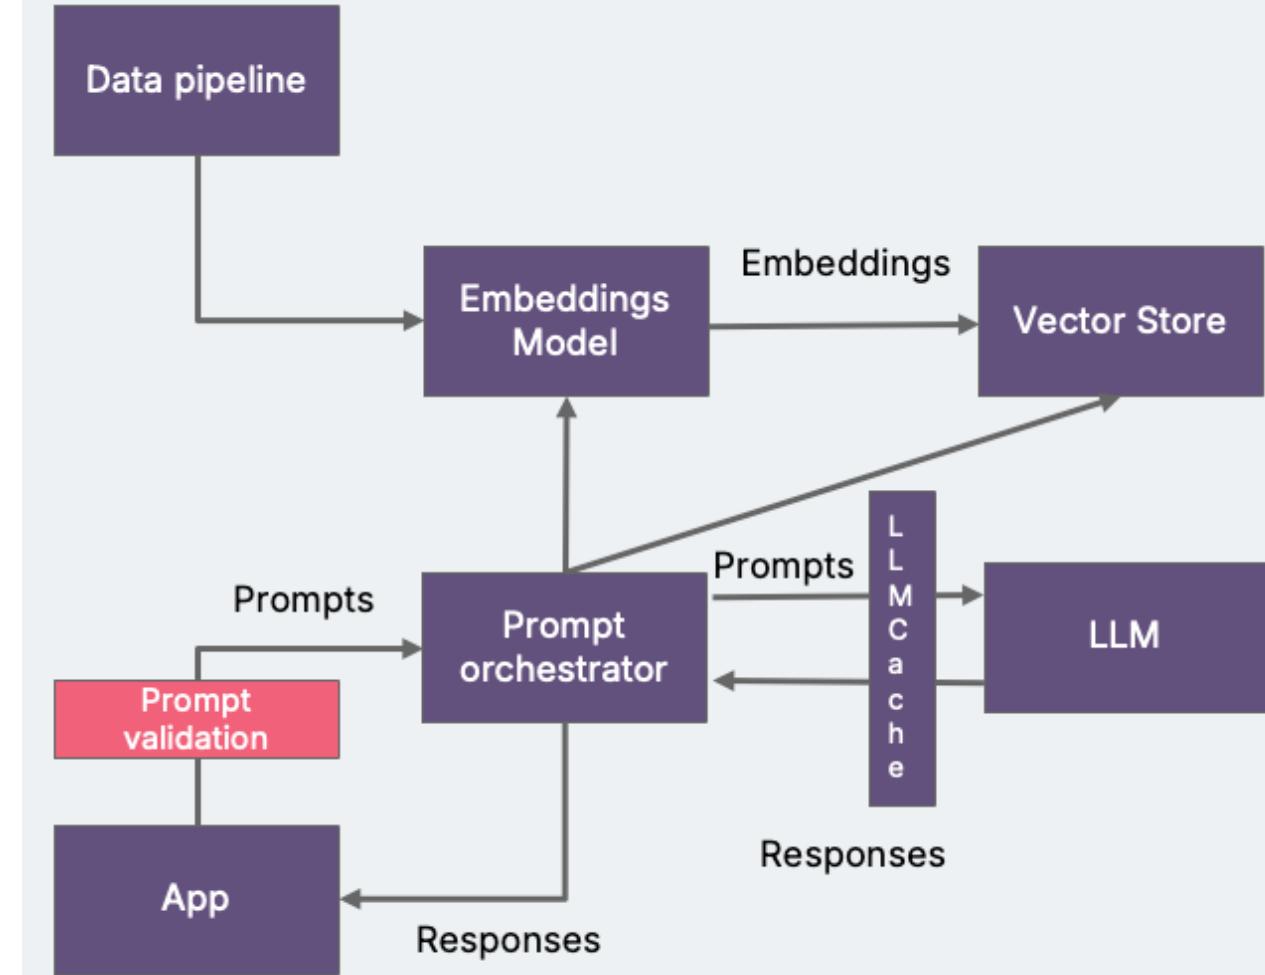


Let's not forget the user interface!

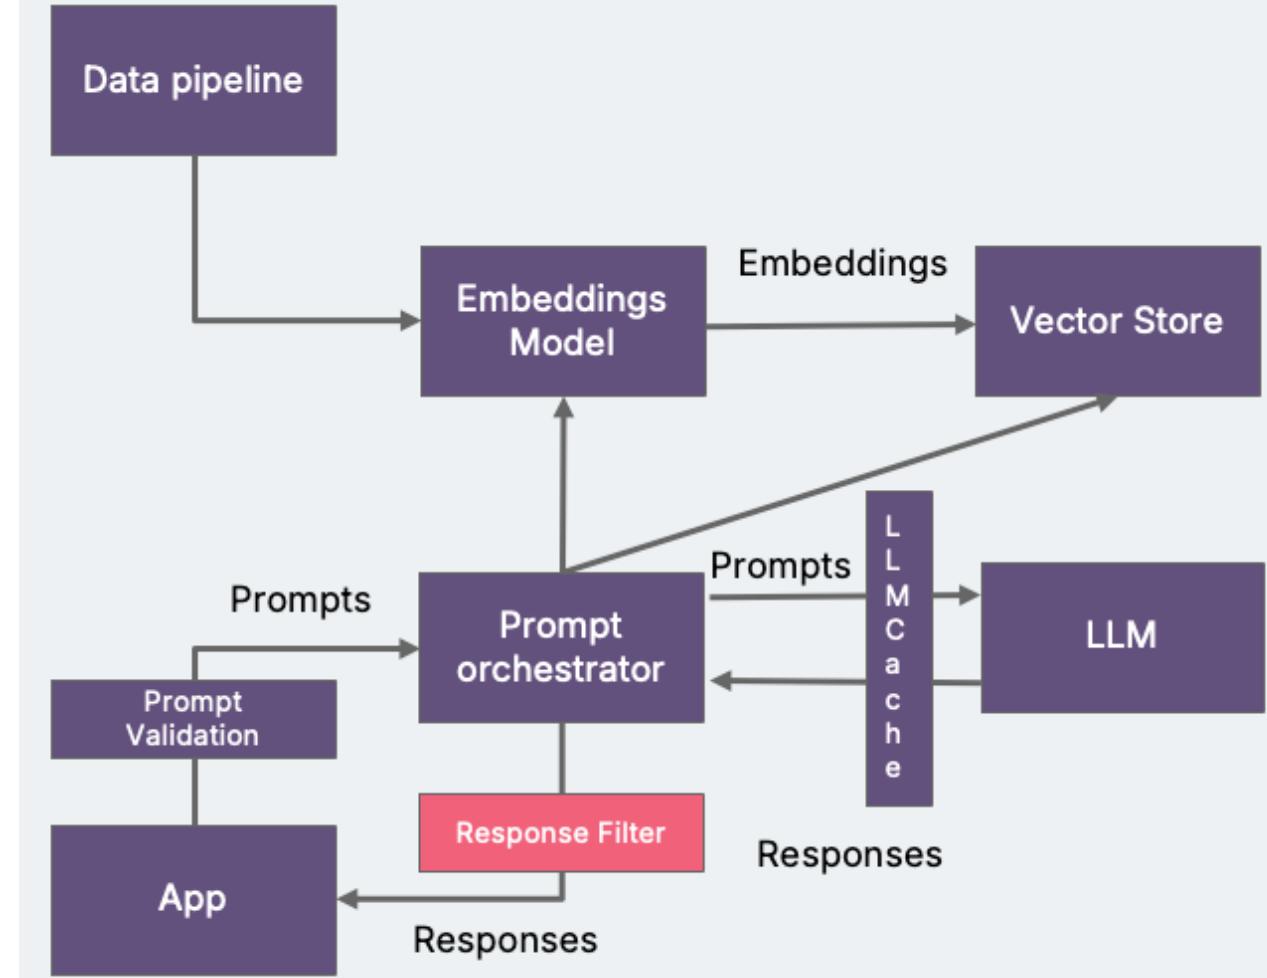


What if our users are untrustworthy?

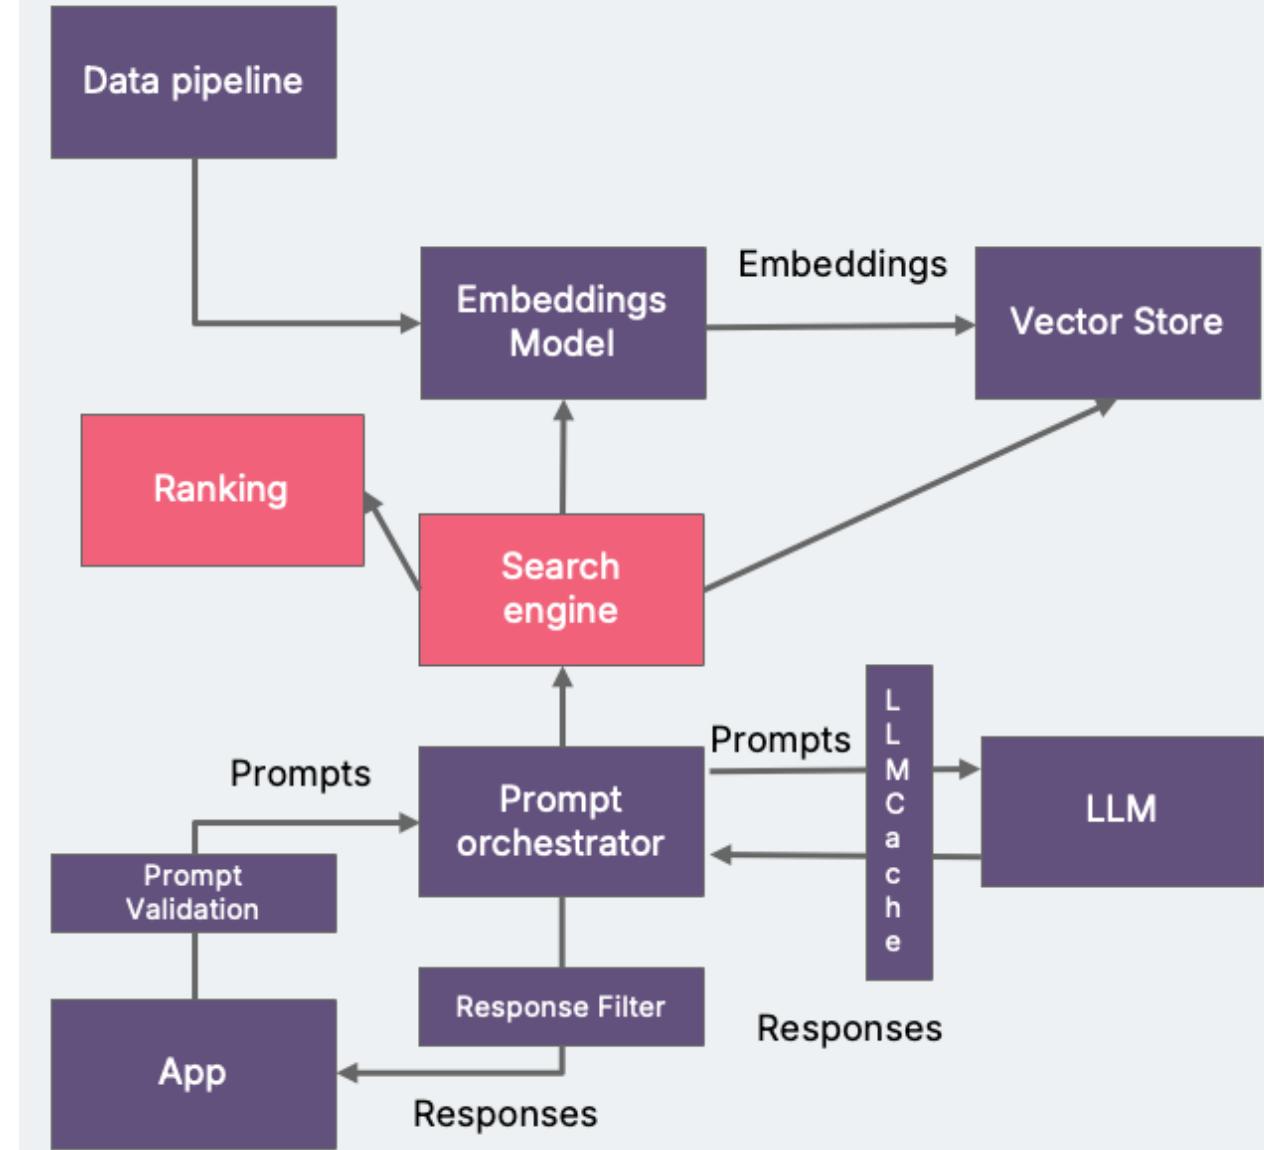
Or vulnerable?



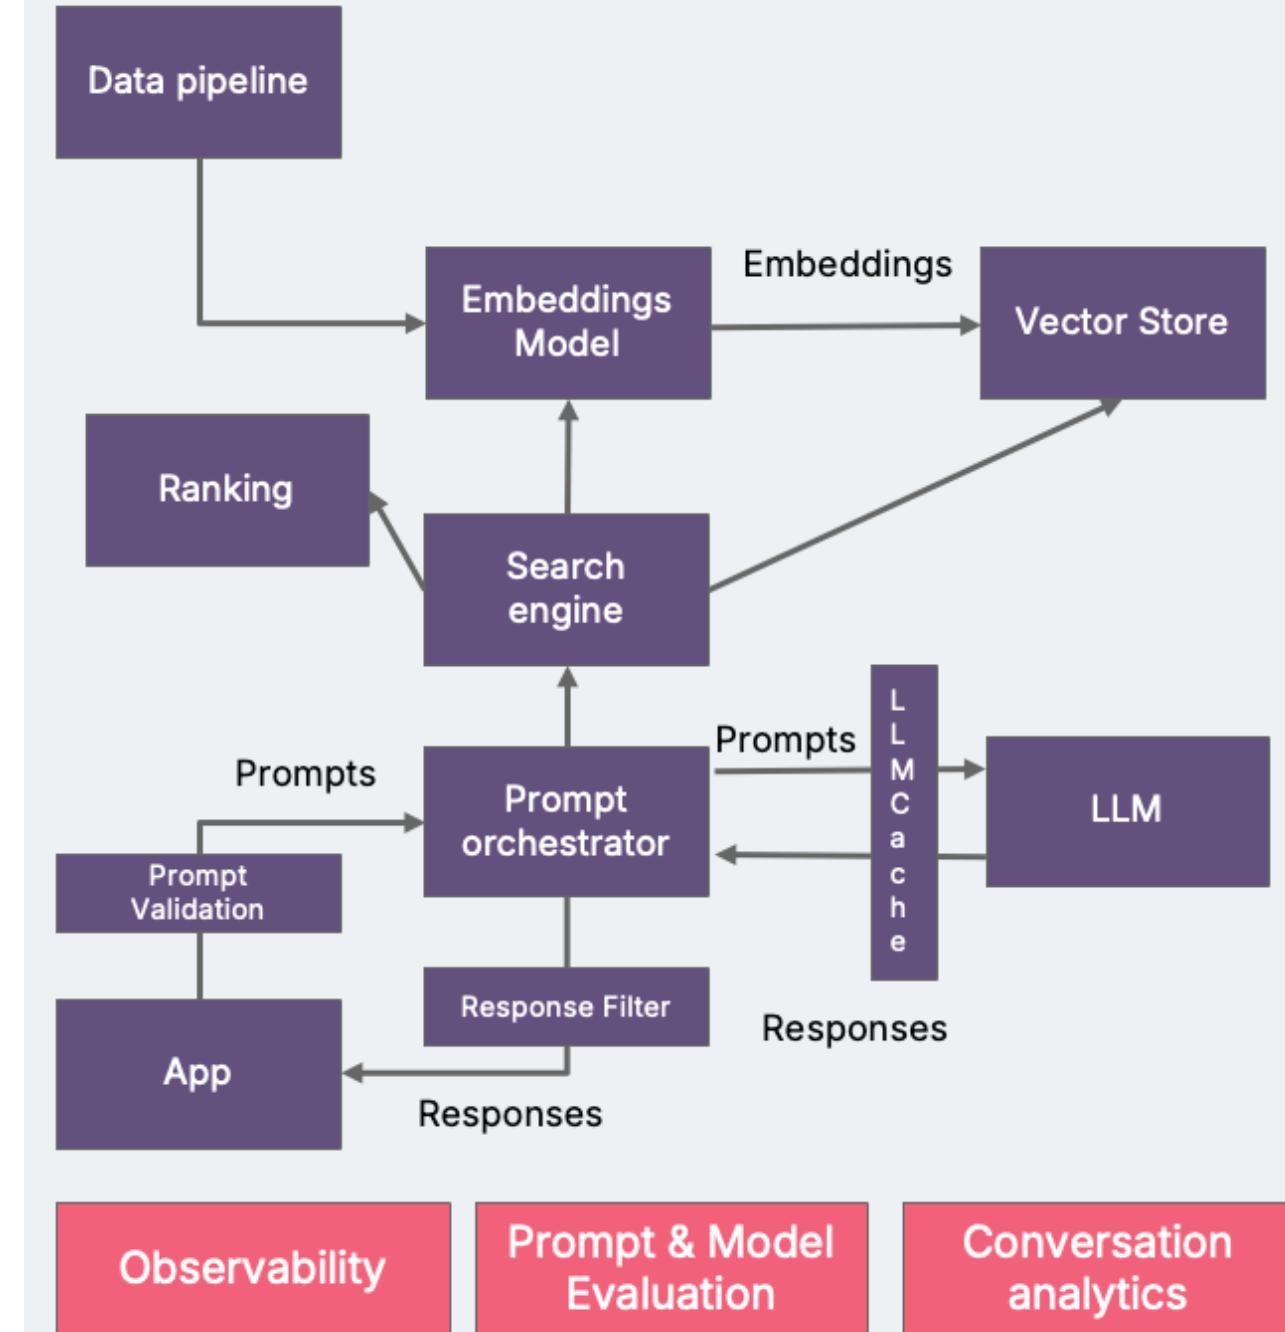
**What if
the response isn't
“appropriate”?**



The context orchestration is not good enough, let's improve the retrieval



We need to manage this!



Deliverables:

- Overview: a *short* narrative describing how the team used AI in the certification system.
- Diagrams: comprehensive and targeted views for each use of AI
- ADRs for AI implementations, including trade-off analysis
- (optional) Pertinent implementation details
- (for semi-final teams) Five-minute video describing the team's approach

Judge's Criteria:

- Innovative use of Generative AI in solution(s)
- Suitability of the solution given the constraints
- Provide appropriate levels of detail
- Use of AI architecture patterns
- Avoidance AI architecture anti-patterns
- Do the architectural characteristics of the additions match the existing architecture?
- Validation and verification of AI results

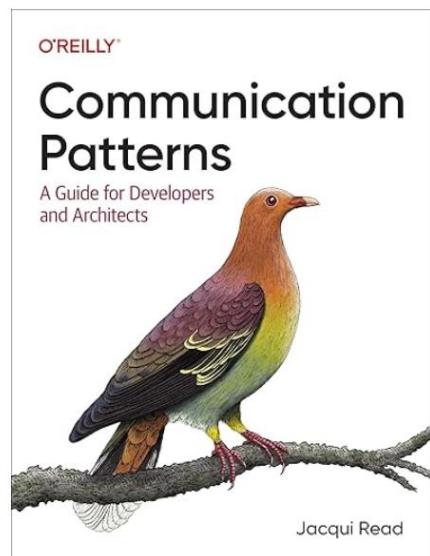


Dates

- All teams must submit this Google Form (<https://forms.gle/JZqZsvoh5u9kNdXX7>) by Thursday, February 6 at 12pm Eastern to participate
- Solutions are due in your GitHub repo by 11:59pm ET Wednesday, February 19
- Semifinalists will be announced at the second event on Wednesday, March 5
- Questions? Email us at katas@oreilly.com

The Katalog

<https://github.com/thekatalog>



www.oreilly.com/library/view/communication-patterns/9781098140533/

The screenshot shows a Mac OS X browser window with the O'Reilly website. The URL bar shows 'oreilly.com'. The main content area features a large image of a person's face, likely a speaker or host. Below the image, the title 'Architectural Katas Winter 2025: AI-Enabled Architecture' is displayed in a large, bold, dark font. To the left of the title is a red 'VIEW ALL EVENTS' button. Above the title, the O'Reilly logo is visible. The top navigation bar includes links for 'TEAMS', 'INDIVIDUALS', 'FEATURES', 'BLOG', 'CONTENT SPONSORSHIP', a search icon, 'SIGN IN', and a 'Try Now' button.

[VIEW ALL EVENTS](#)

Architectural Katas Winter 2025: AI-Enabled Architecture

Published by [O'Reilly Media, Inc.](#)

Intermediate

Challenge your architecture knowledge and your creativity in leveraging generative AI

[What you'll learn](#) [Is this live event for you?](#) [Schedule](#)

By registering for Architectural Katas, you get access to the three live online events and can vote for the winning presentation. Registration does not mean you're also participating in the Architectural Katas challenge. If you'd like to take part, you must gather a team of three to five people and submit your team name to O'Reilly. Team registration will open February 5 following the first event. Your team must include at least one O'Reilly member in order to participate in the Architectural Katas challenge and attend the live events. The first 100 teams to sign up will be selected to participate.

About Architectural Katas: What could be better than practicing software architecture with your handpicked dream team of architects? This year, the Katas challenge will test both your architecture knowledge and your creativity in leveraging generative AI to help with your solution. Come to participate or just to watch, learn, and get ideas on how to bring these tools into your architecture.

Participants in Architectural Katas, moderated by Neal Ford and Mark Richards, work together in small groups of their choosing to solve a challenge a fictional organization is facing. Everyone gets the chance to practice their craft with a real software architecture problem in a safe and social environment, and the finalists and winners earn well-deserved bragging rights (not to mention an intriguing line on their résumé). Plus, this year attendees will get to see how some of the brightest minds are using generative AI to enhance their work.

How it works: You put together a first-rate team of three to five people, ready to tackle an architecture challenge. We'll share the architecture problem with you at the kickoff on February 5. Then your team will have to solve it, working in whatever way is best for you (video calls, group chat, shared docs, etc.). **Registration opens February 5 following the first event; the first 100 teams to sign up will be selected to participate.**

Not ready to compete but want to be part of the action? Register for the event and join us to see how Architectural Katas works, cast your vote for the winning team, and learn how to successfully present architecture plans to stakeholders.

What you'll learn and how you can apply it

- Identify and design architectural components
- Explore techniques to analyze and select the right architecture patterns
- Learn how to diagram and document an architectural solution
- Develop better communication skills to champion architecture decisions to developers and stakeholders

Feb 5

11am-1pm Eastern Standard Time

March 5

11am-1pm Eastern Standard Time

March 27

11am-1pm Eastern Daylight Time

93 Spots Remaining

[Sign up for a free trial!](#)

or [sign in](#)

Your Hosts



Neal Ford

Neal Ford is a director, software architect, and meme wrangler at Thoughtworks, a software company and a community of passionate, purpose-led individuals who think disruptively to deliver technology to address the...

[Read more](#)

X Ø Q



Mark Richards

Mark Richards is an experienced, hands-on software architect involved in the architecture, design, and implementation of microservices architectures and other distributed systems in a variety of technologies. He has been in the...

[Read more](#)

in X Ø Q

Associated roles

Cloud solutions architect

Data architect

ML/AI architect

Senior software developer

+2 more

Skill covered

Software Architecture

Architecture Katas

Winter 2025: Architecture & AI

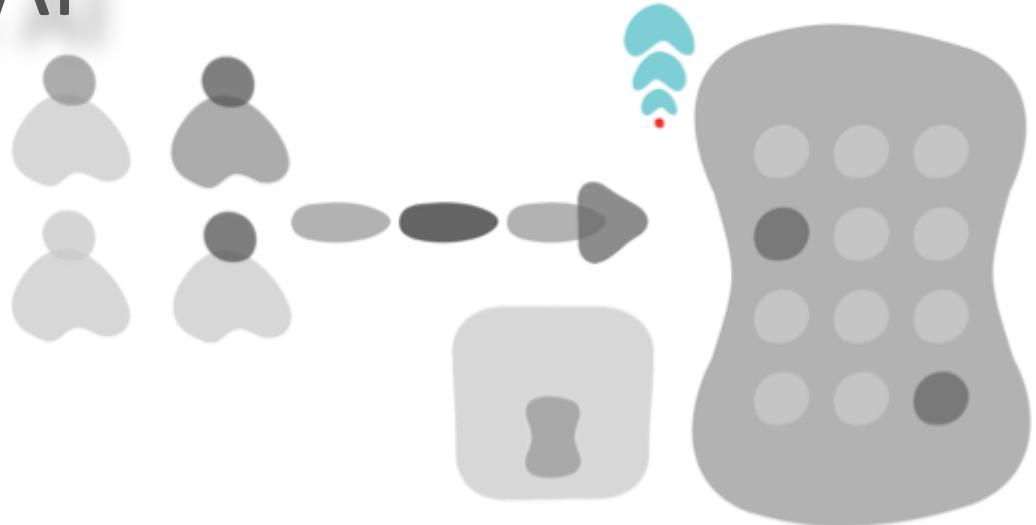


Neal Ford

Thoughtworks

Director / Software Architect / Meme Wrangler

<https://www.nealford.com>



Mark Richards

Independent Consultant

Hands-on Software Architect, Published Author

Founder, DeveloperToArchitect.com

@markrichardssa

Contest Kickoff